

16/06/2025

# Savvy

## AI - Powered Personal Finance Management Application

Graduation Project II - A.Y. 2024-2025

### Prepared By:

Basmala Salama Hassan

Elham Hamed Mahmoud

Irinie Magued Sabry

Menna Allah Samy

Menna Allah Saeed

### Supervised By:

Dr. Reem Essameldin





# Savvy

AI - Powered Personal Finance Management  
Application



**SAVVY**

# Overview

- **Key Features Overview**

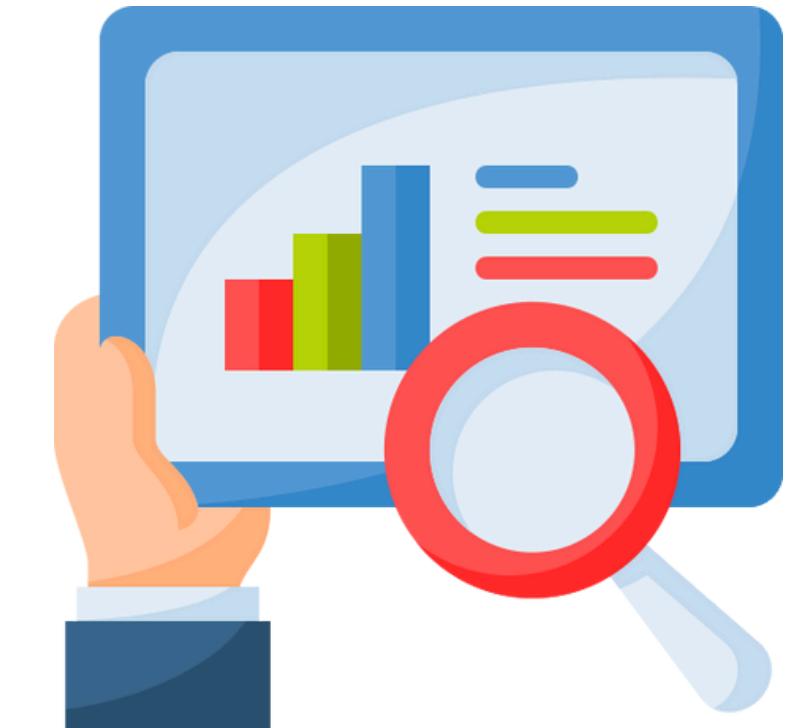
- Automated Transaction Tracking
- Budget Planning
- Goal Management
- Chatbot & Spending Forecast (AI Features)

- **Tech Stack:**

- Frontend: Flutter
- Backend: FastAPI + Supabase
- ML Models: LSTM & LangChain RAG & Transactions chatbot

- **Project Objectives**

- Help users track income and expenses
- Provide smart financial insights using AI
- Enable personalized goal setting and budgeting
- Predict future spending trends





# **Savvy**

## **Backend Development**



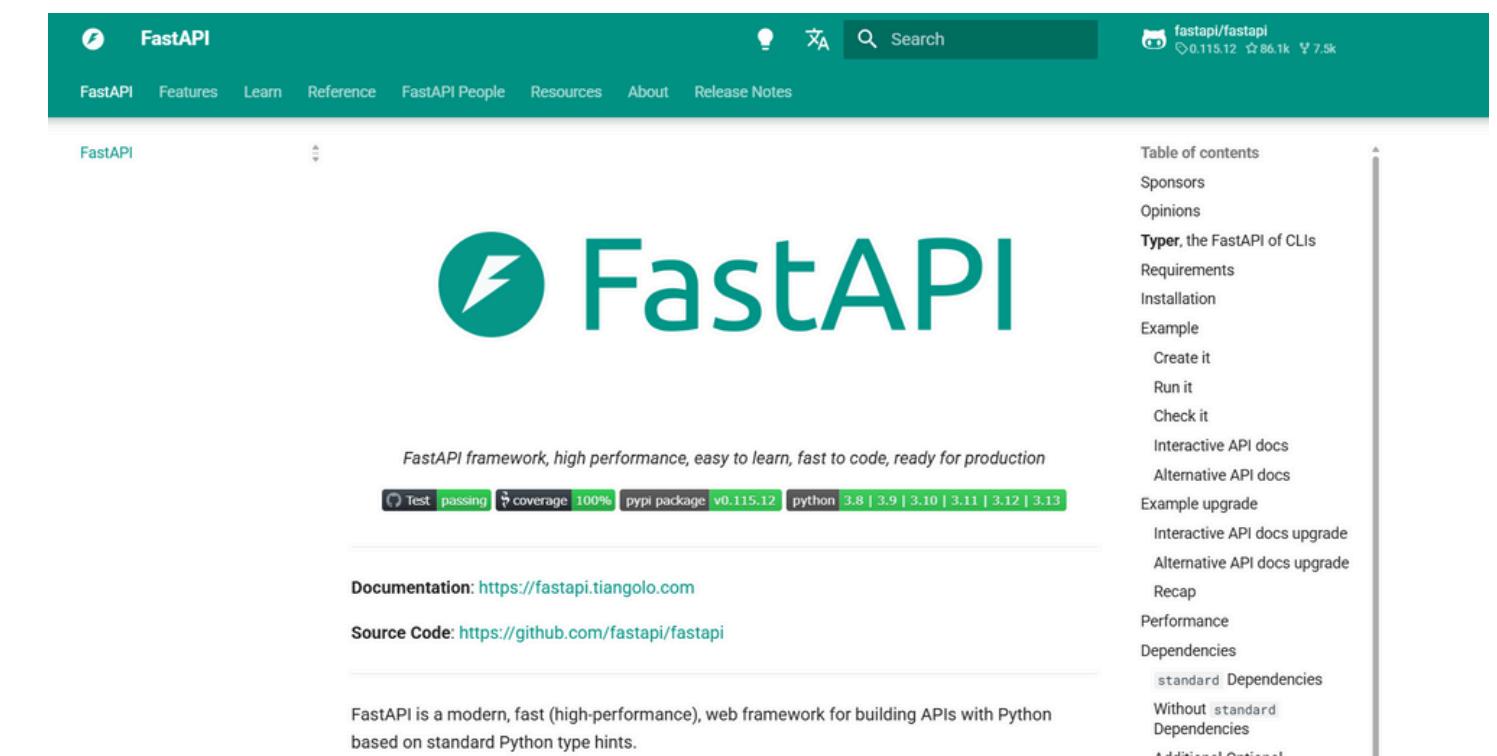
# Backend Development



## FastAPI

FastAPI is a modern web framework that is relatively fast and used for building APIs with Python

- High performance using asynchronous Python
- Automatic documentation with Swagger & ReDoc
- Type-safety with Pydantic
- Fast development cycle



The screenshot shows the official FastAPI website. At the top, there's a navigation bar with links for FastAPI, Features, Learn, Reference, FastAPI People, Resources, About, and Release Notes. Below the navigation is a large teal header with the FastAPI logo (a stylized lightning bolt) and the text "FastAPI". To the right of the logo, there's a search bar and a GitHub icon. The main content area features a large "FastAPI" title with a lightning bolt icon. Below it, there's a brief description: "FastAPI framework, high performance, easy to learn, fast to code, ready for production". It includes links for "Documentation: <https://fastapi.tiangolo.com>" and "Source Code: <https://github.com/fastapi/fastapi>". On the far right, there's a sidebar with a table of contents and various links related to the FastAPI framework.

# What is Supabase and Why We Chose It?



Supabase is an open-source Backend-as-a-Service (BaaS) platform that provides a full backend solution powered by PostgreSQL – the world's most advanced open-source relational database.



## Row-Level Security (RLS)

Ensures secure, per-user data access using Row-Level Security—so users only see their own data.



## Real-Time Updates

Enables dynamic updates to transactions and budgets without page refreshes, improving UX.



## SQL Flexibility

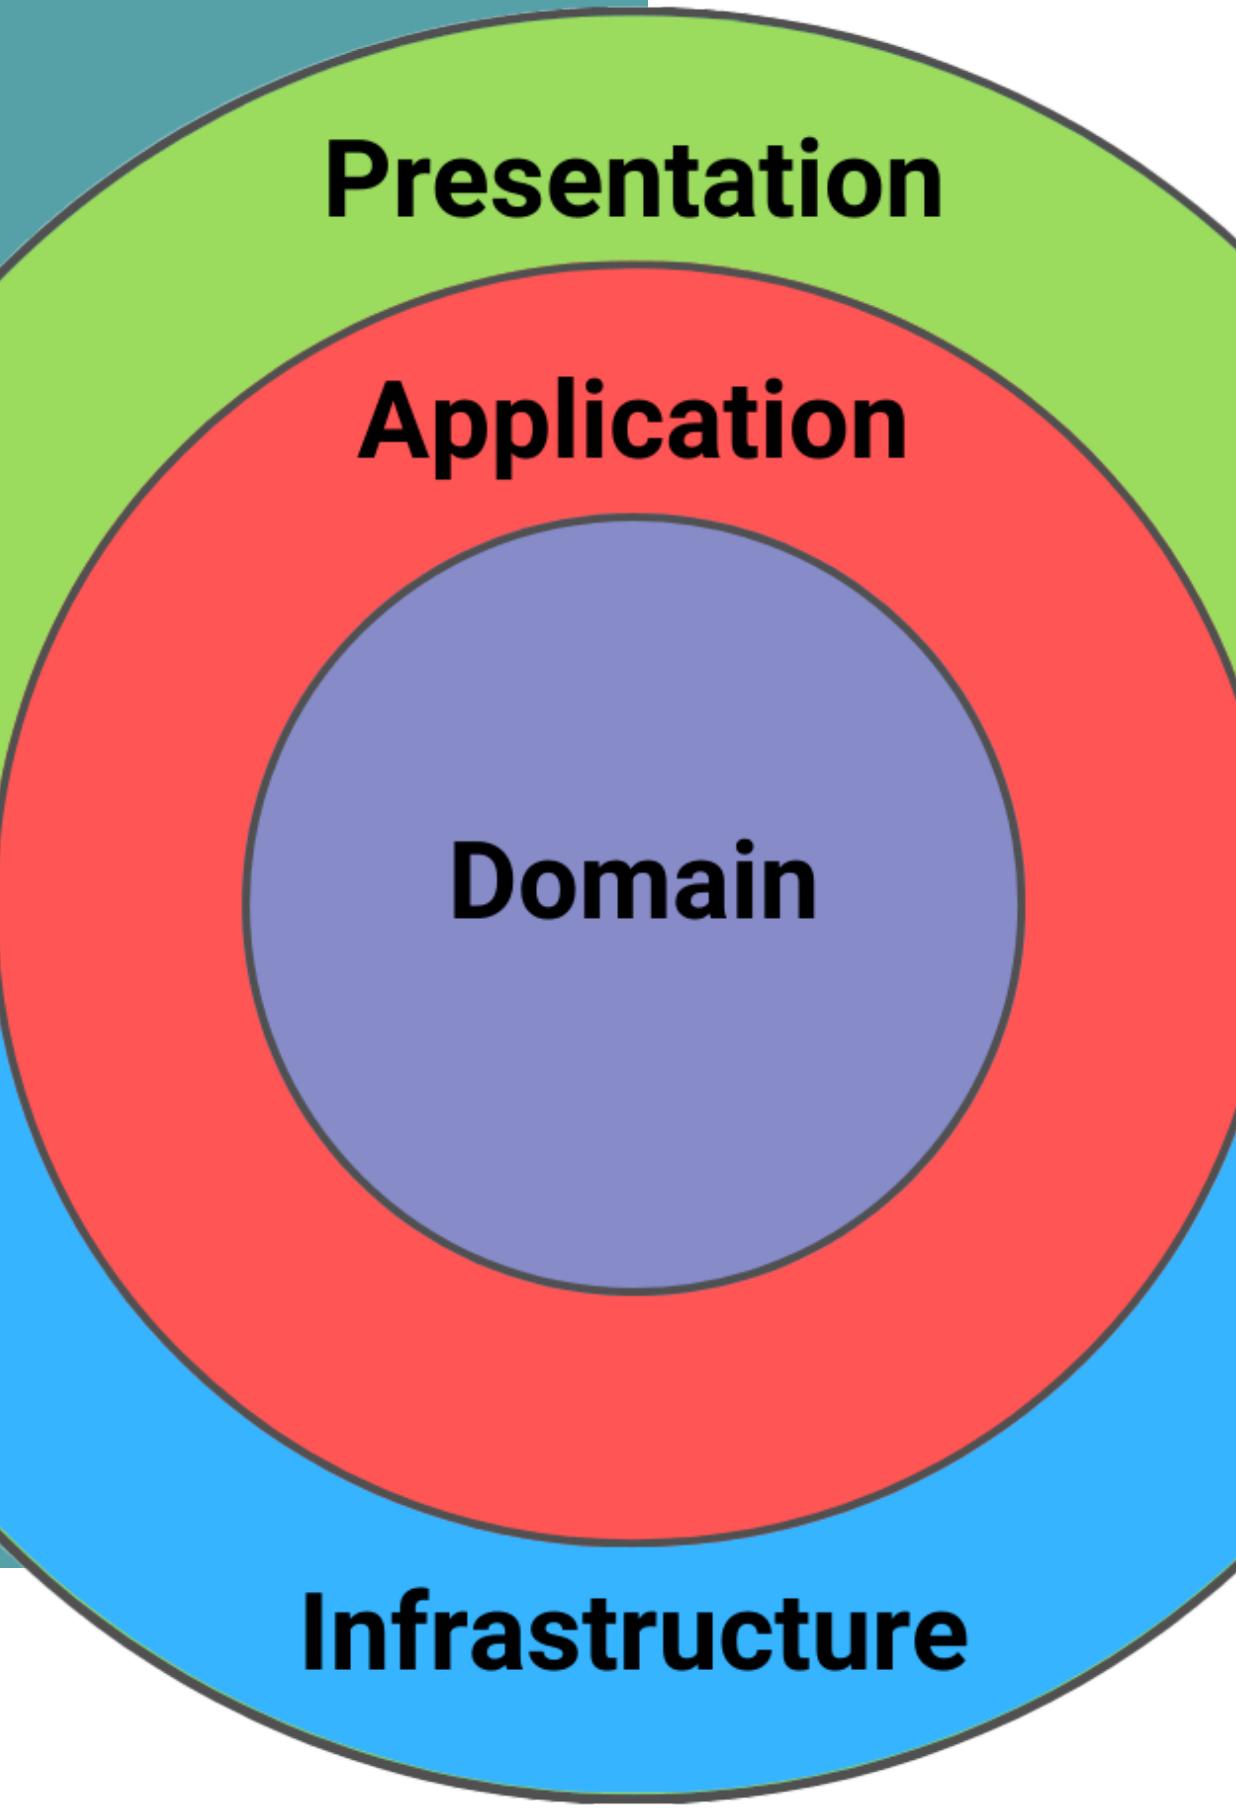
Advanced querying, full-text search, and time-series support helped us build powerful analytics and prediction features.



## User Authentication & Authorization

- Fully integrated sign-up / login / session handling
- JWT-based authentication with Row-Level Security (RLS) to enforce per-user access to data.

# Clean Architecture



Clean Architecture is a software design approach that emphasizes long-term modularity, and testability by enforcing a strict separation of concerns, particularly **between core business logic and external systems**

**Presentation Layer**  
(API endpoints)

**Service Layer**  
(Business logic)

**Infrastructure Layer**  
(Data access)

**Domain Layer**  
(Entities, Models)

# **Repository Design Pattern**

A structural design pattern that

- Abstracts how data is fetched, saved, updated, or deleted.
- Acts like a controlled in-memory collection for domain entities.
- Hides the details of data persistence from the application logic.
- Promotes modular, testable, and maintainable code.

**What app does VS How data stored**

# Frontend

The Flutter logo consists of three overlapping blue rectangles of increasing transparency from top-left to bottom-right. The word "Flutter" is centered in bold black font.

# Flutter

# Flutter

## App Development



Google

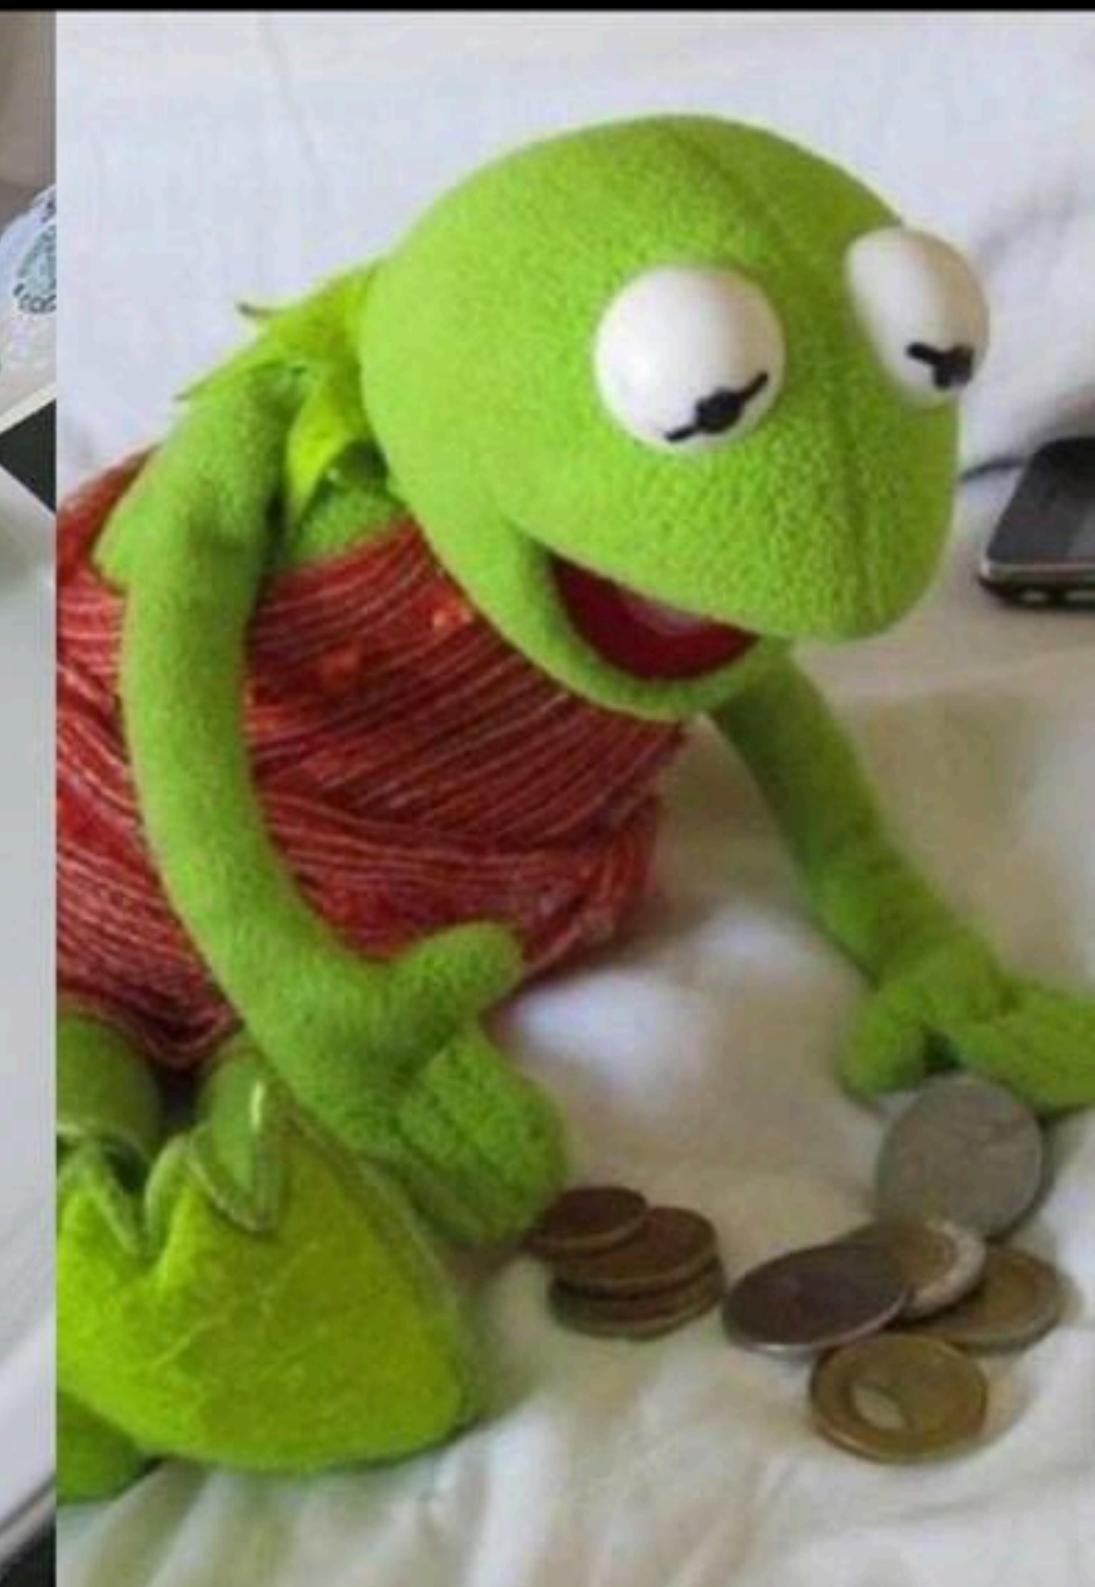
**Where is my money??!**

**When will I reach my goal?**



**Beginning of  
the month**

**End of  
the month**



# Demo Time



9:03

100%

## Add Expense

Date

15/06/2025

Food

Amount

EGP 00.00

Description (optional)

Save



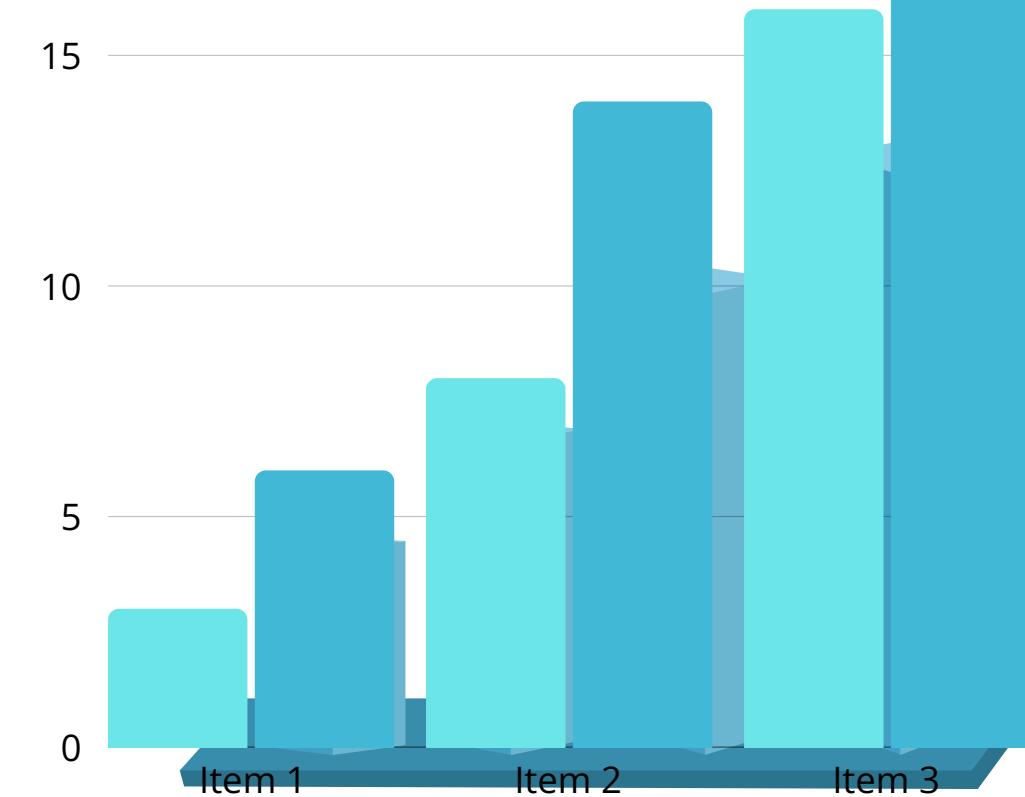
0 1 2 3 4 5 6 7 8 9 0

# Analytics

`fl_chart`

We made sure the analytics update in real time  
by using **state management with Riverpod**.

Any time the user adds a transaction or saves  
money toward a goal, our providers  
automatically recalculate the totals and trigger  
a UI update.



9:12

69%

## Financial Analytics

Daily

Weekly

Monthly

Net Savings

**EGP 67710.00**

Income

**EGP 111000.00**

Expense

**EGP 43290.00**

85200

50000

0

Wk 5/19

Wk 5/26

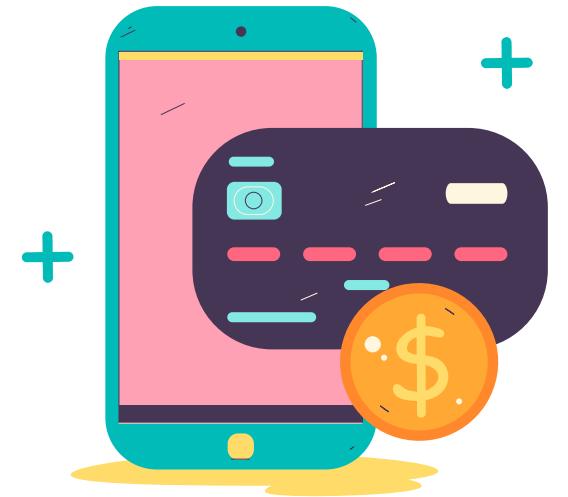
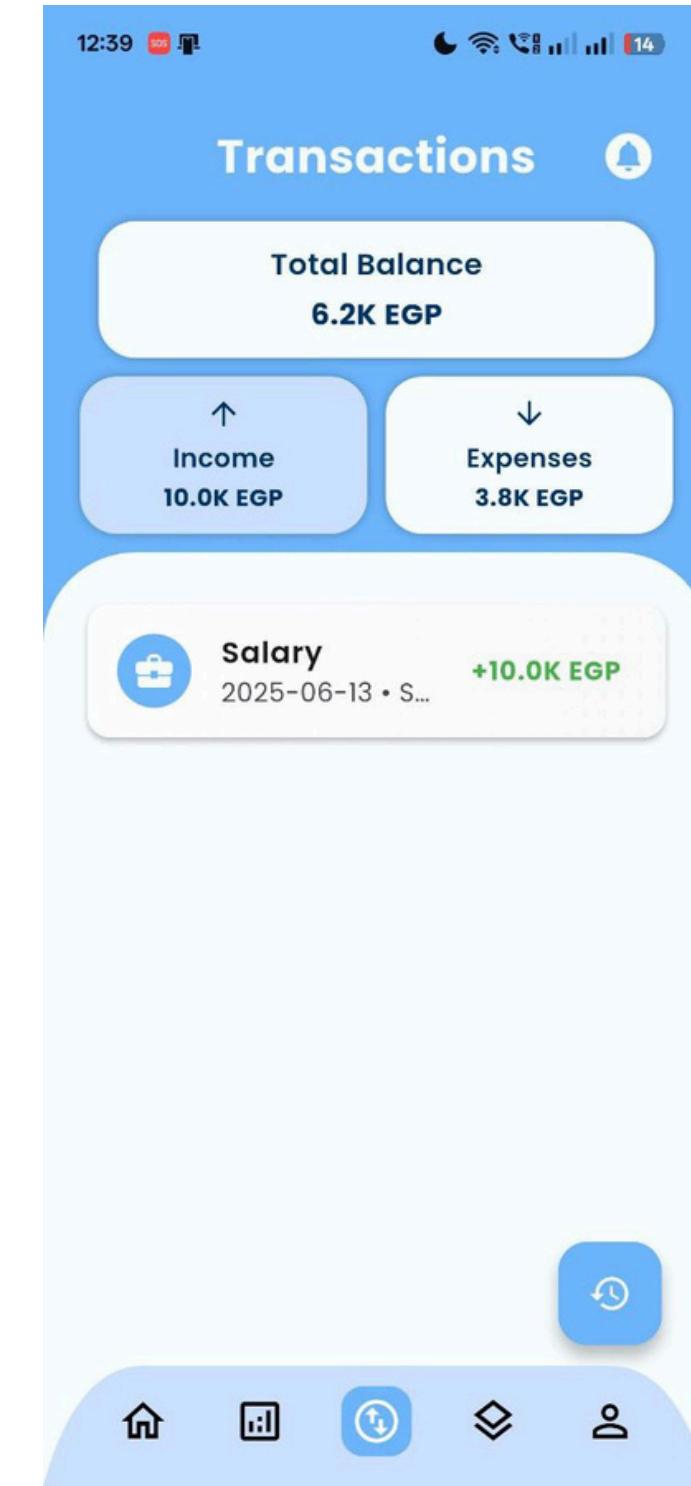
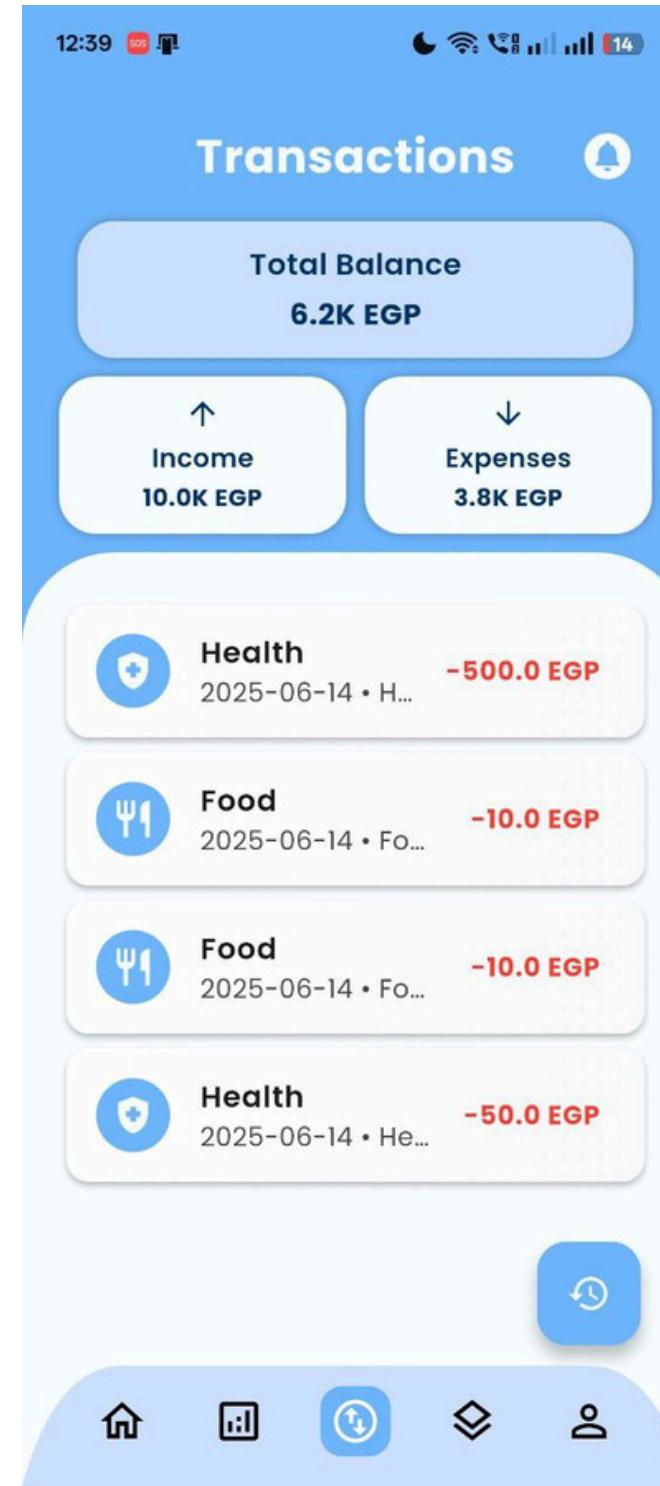
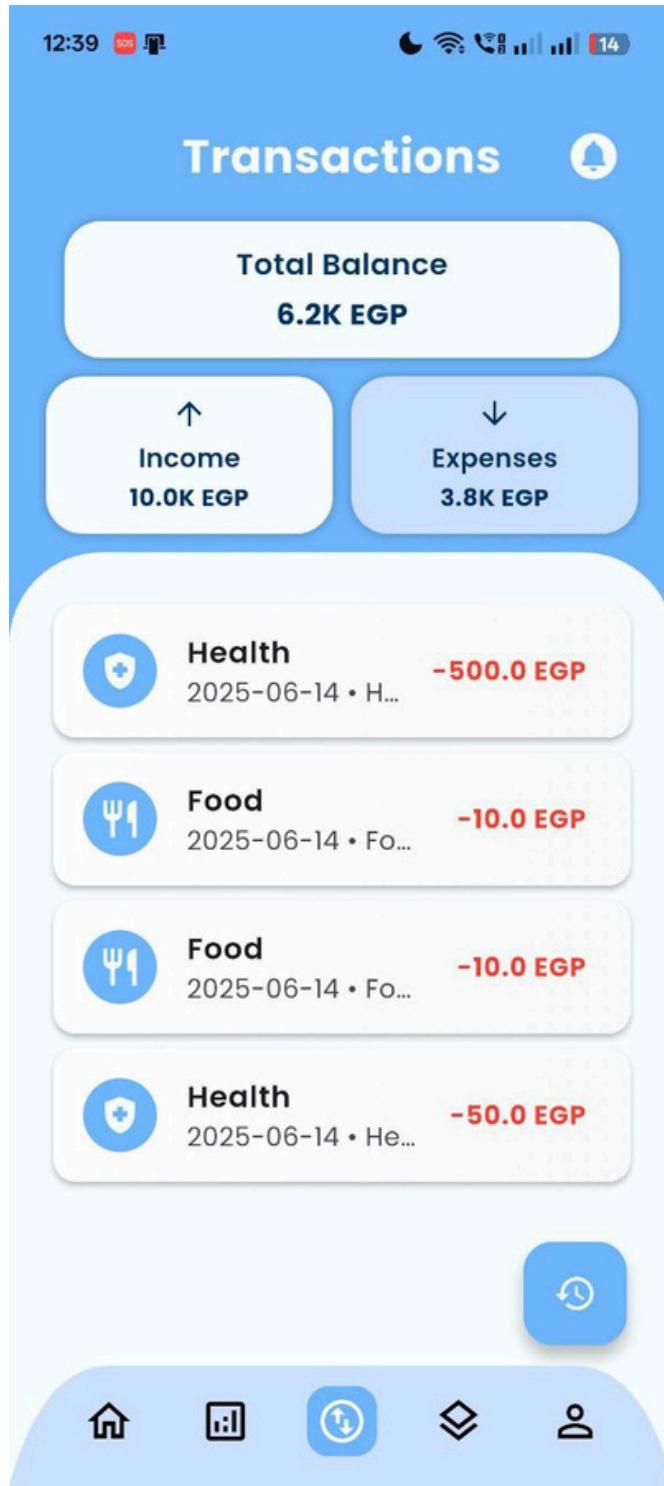
Wk 6/2

Wk 6/9

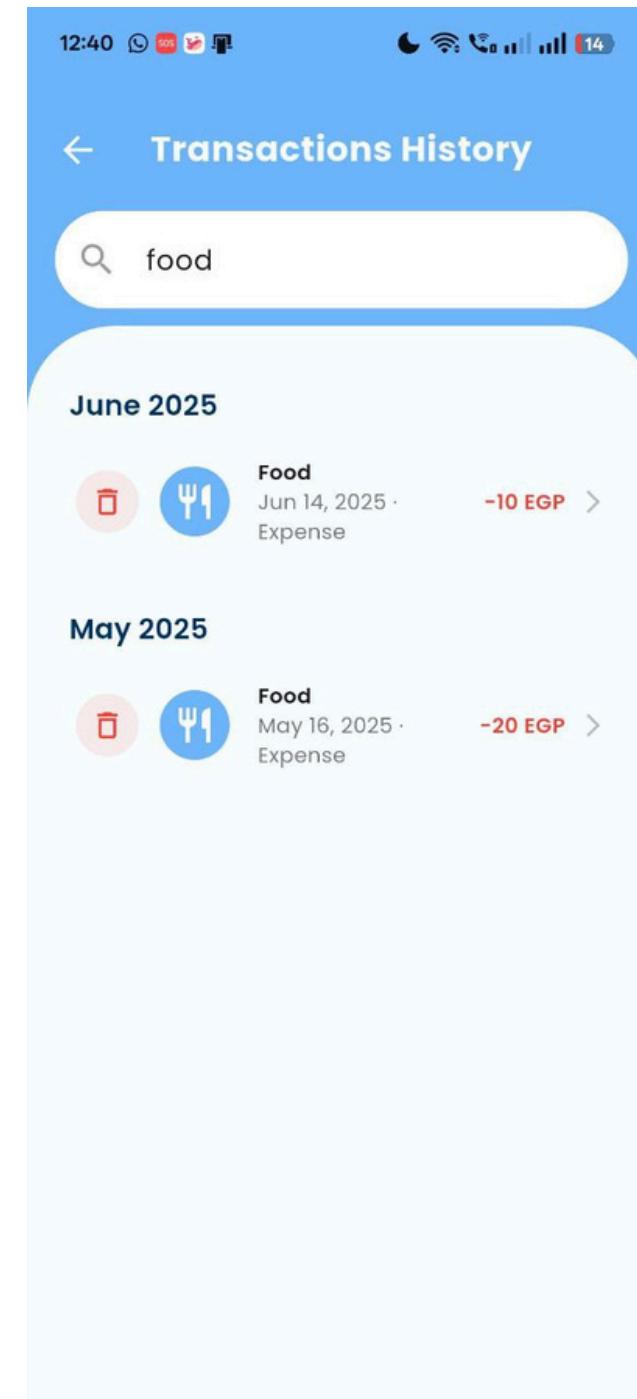
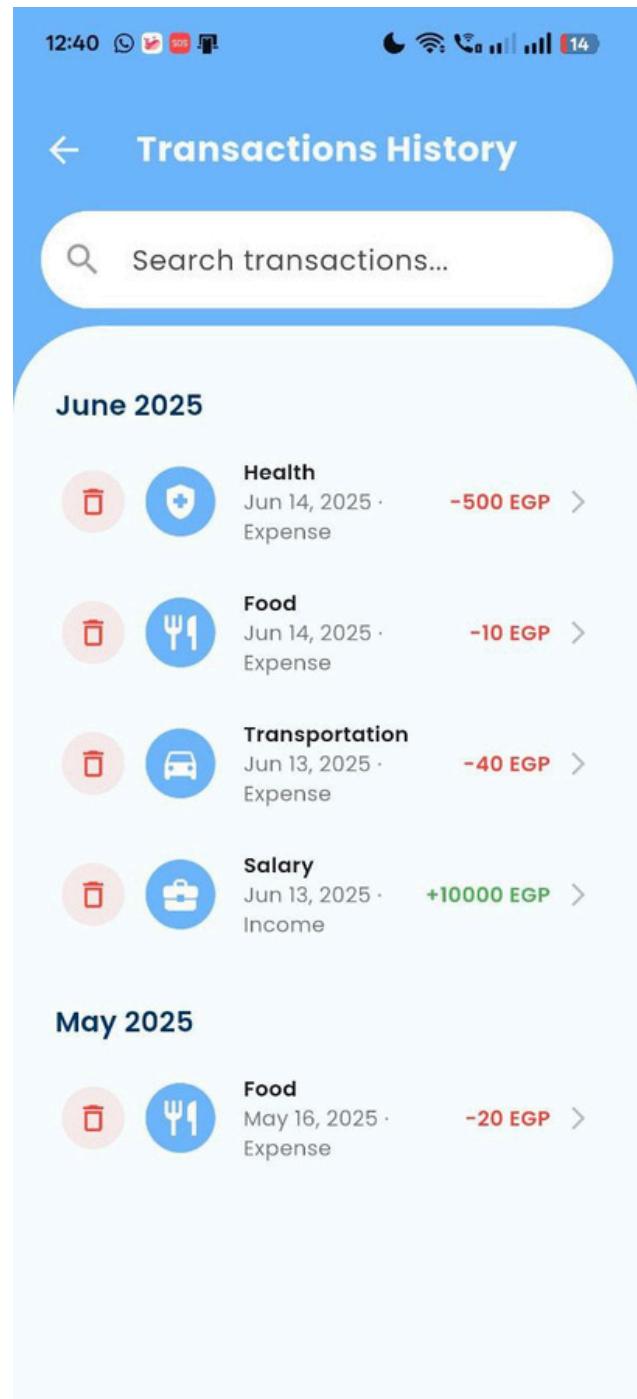
### Savings Goals



# Transactions



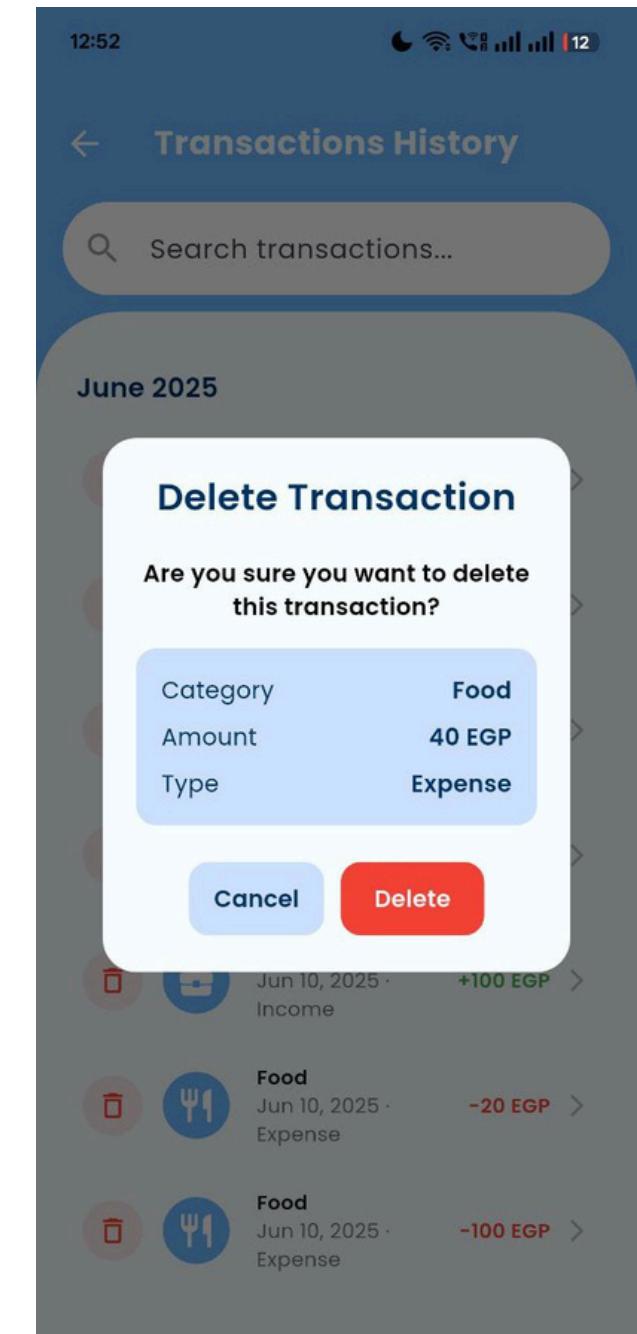
# Transactions history



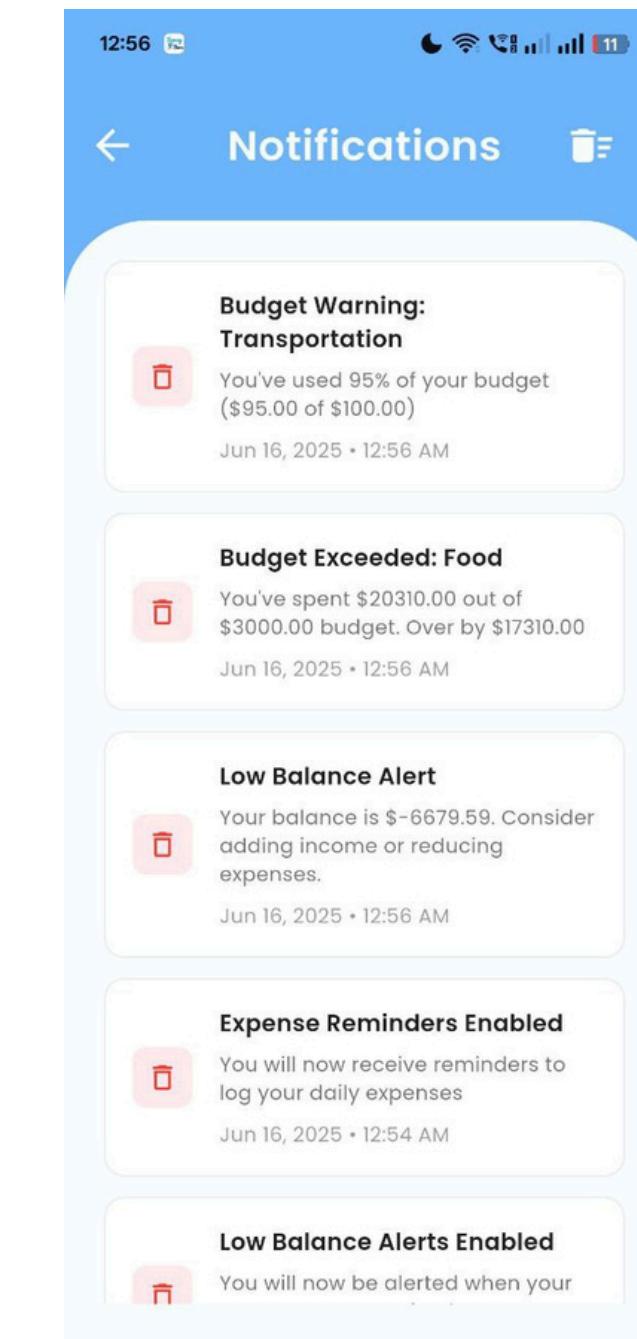
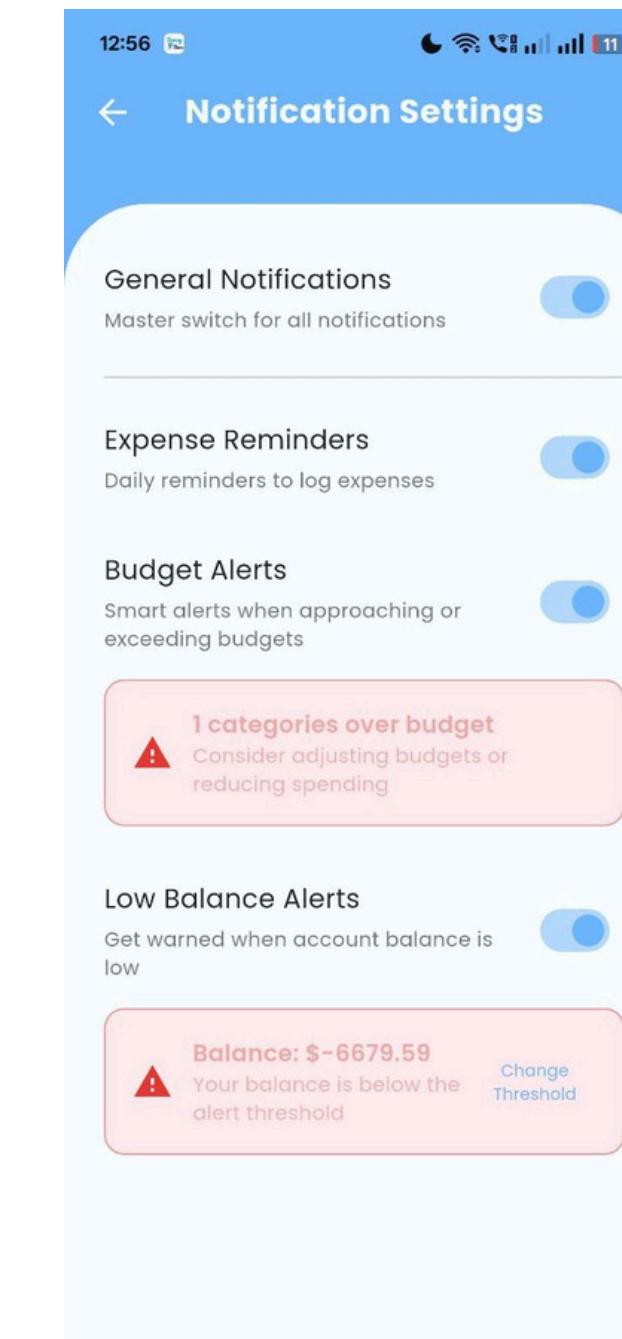
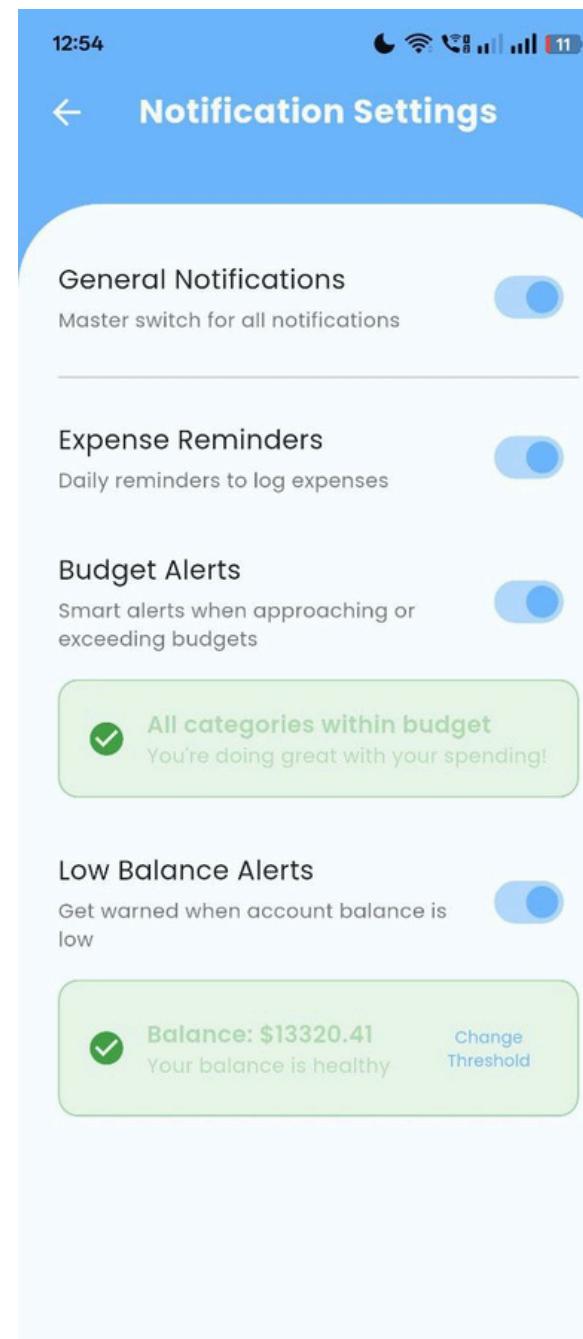
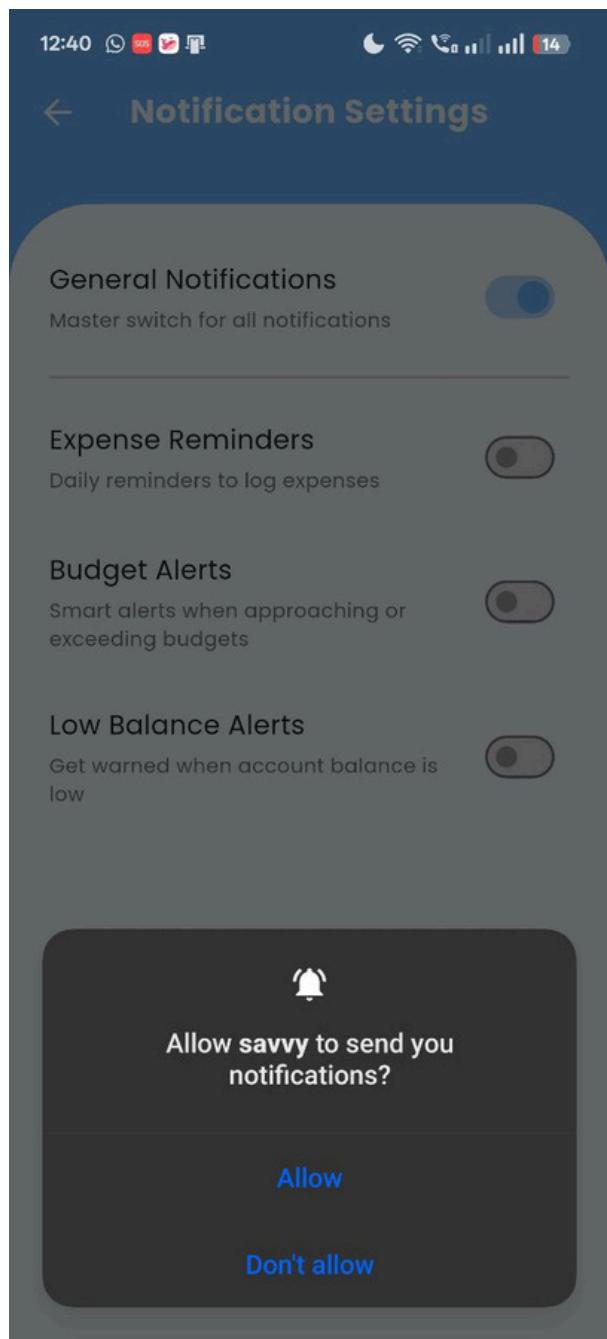
This screenshot shows the Transaction Details screen. The transaction being edited is an expense for Food on Jun 11, 2025, with an amount of EGP 40. The form fields include:

- Transaction Type: Expense
- Category: Food
- Amount: EGP 40
- Description: Enter description (optional)
- Date: Jun 11, 2025

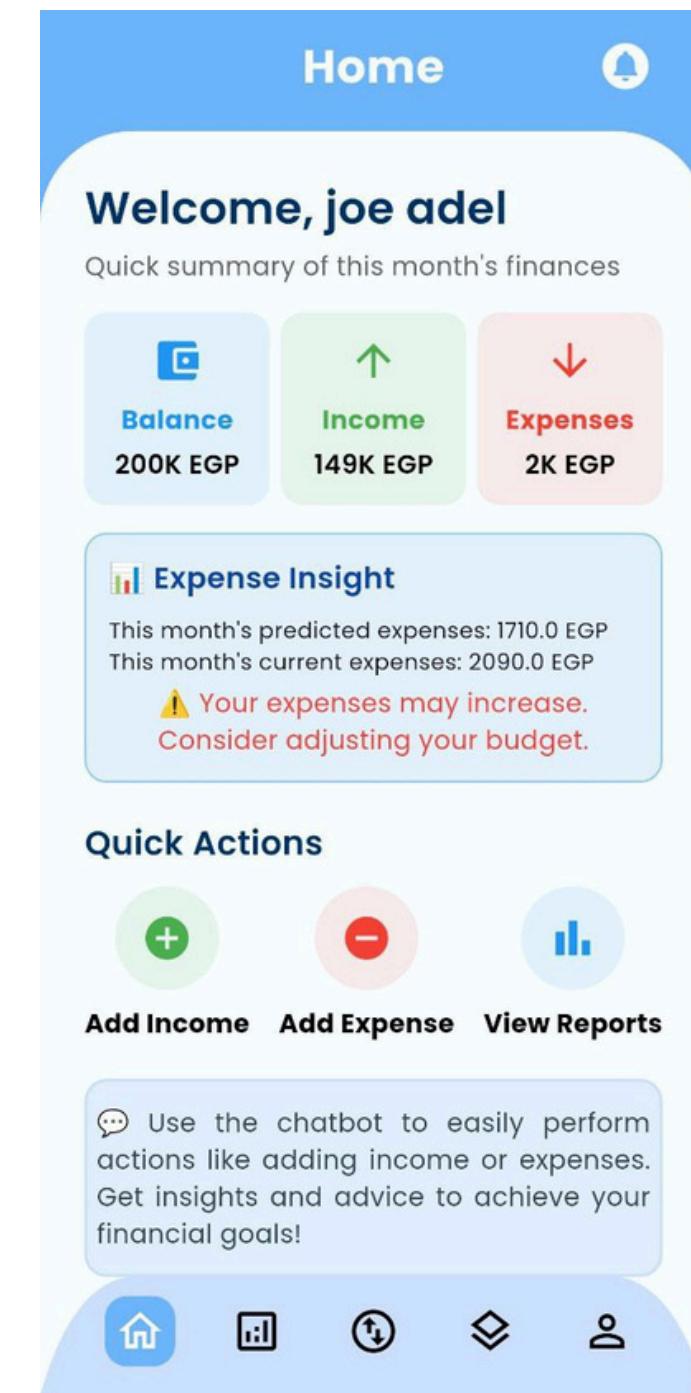
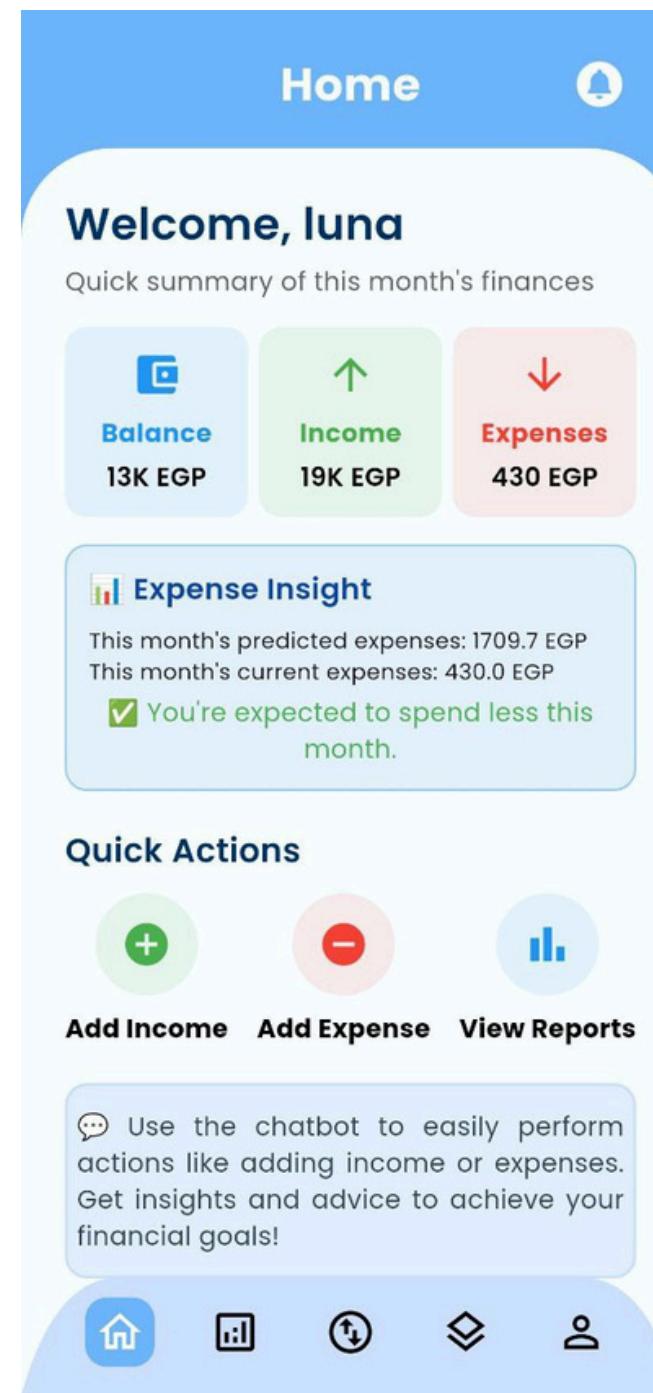
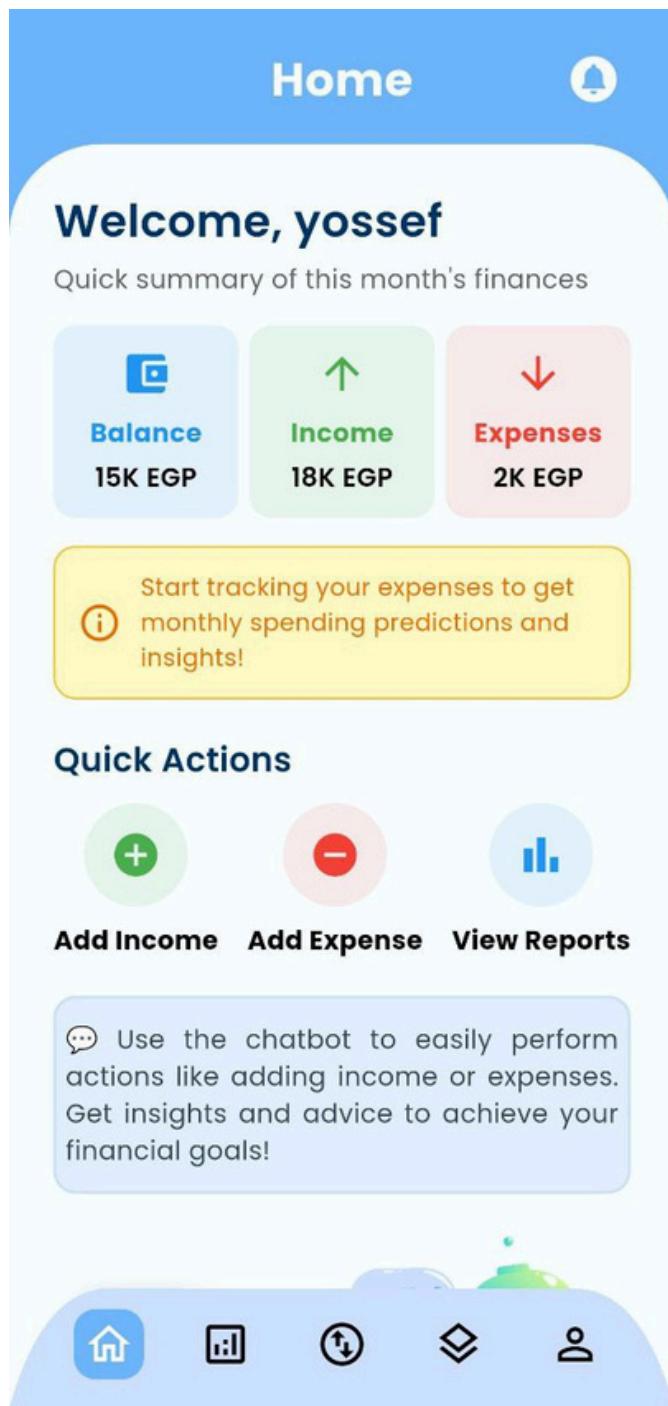
A "Save Changes" button is at the bottom.



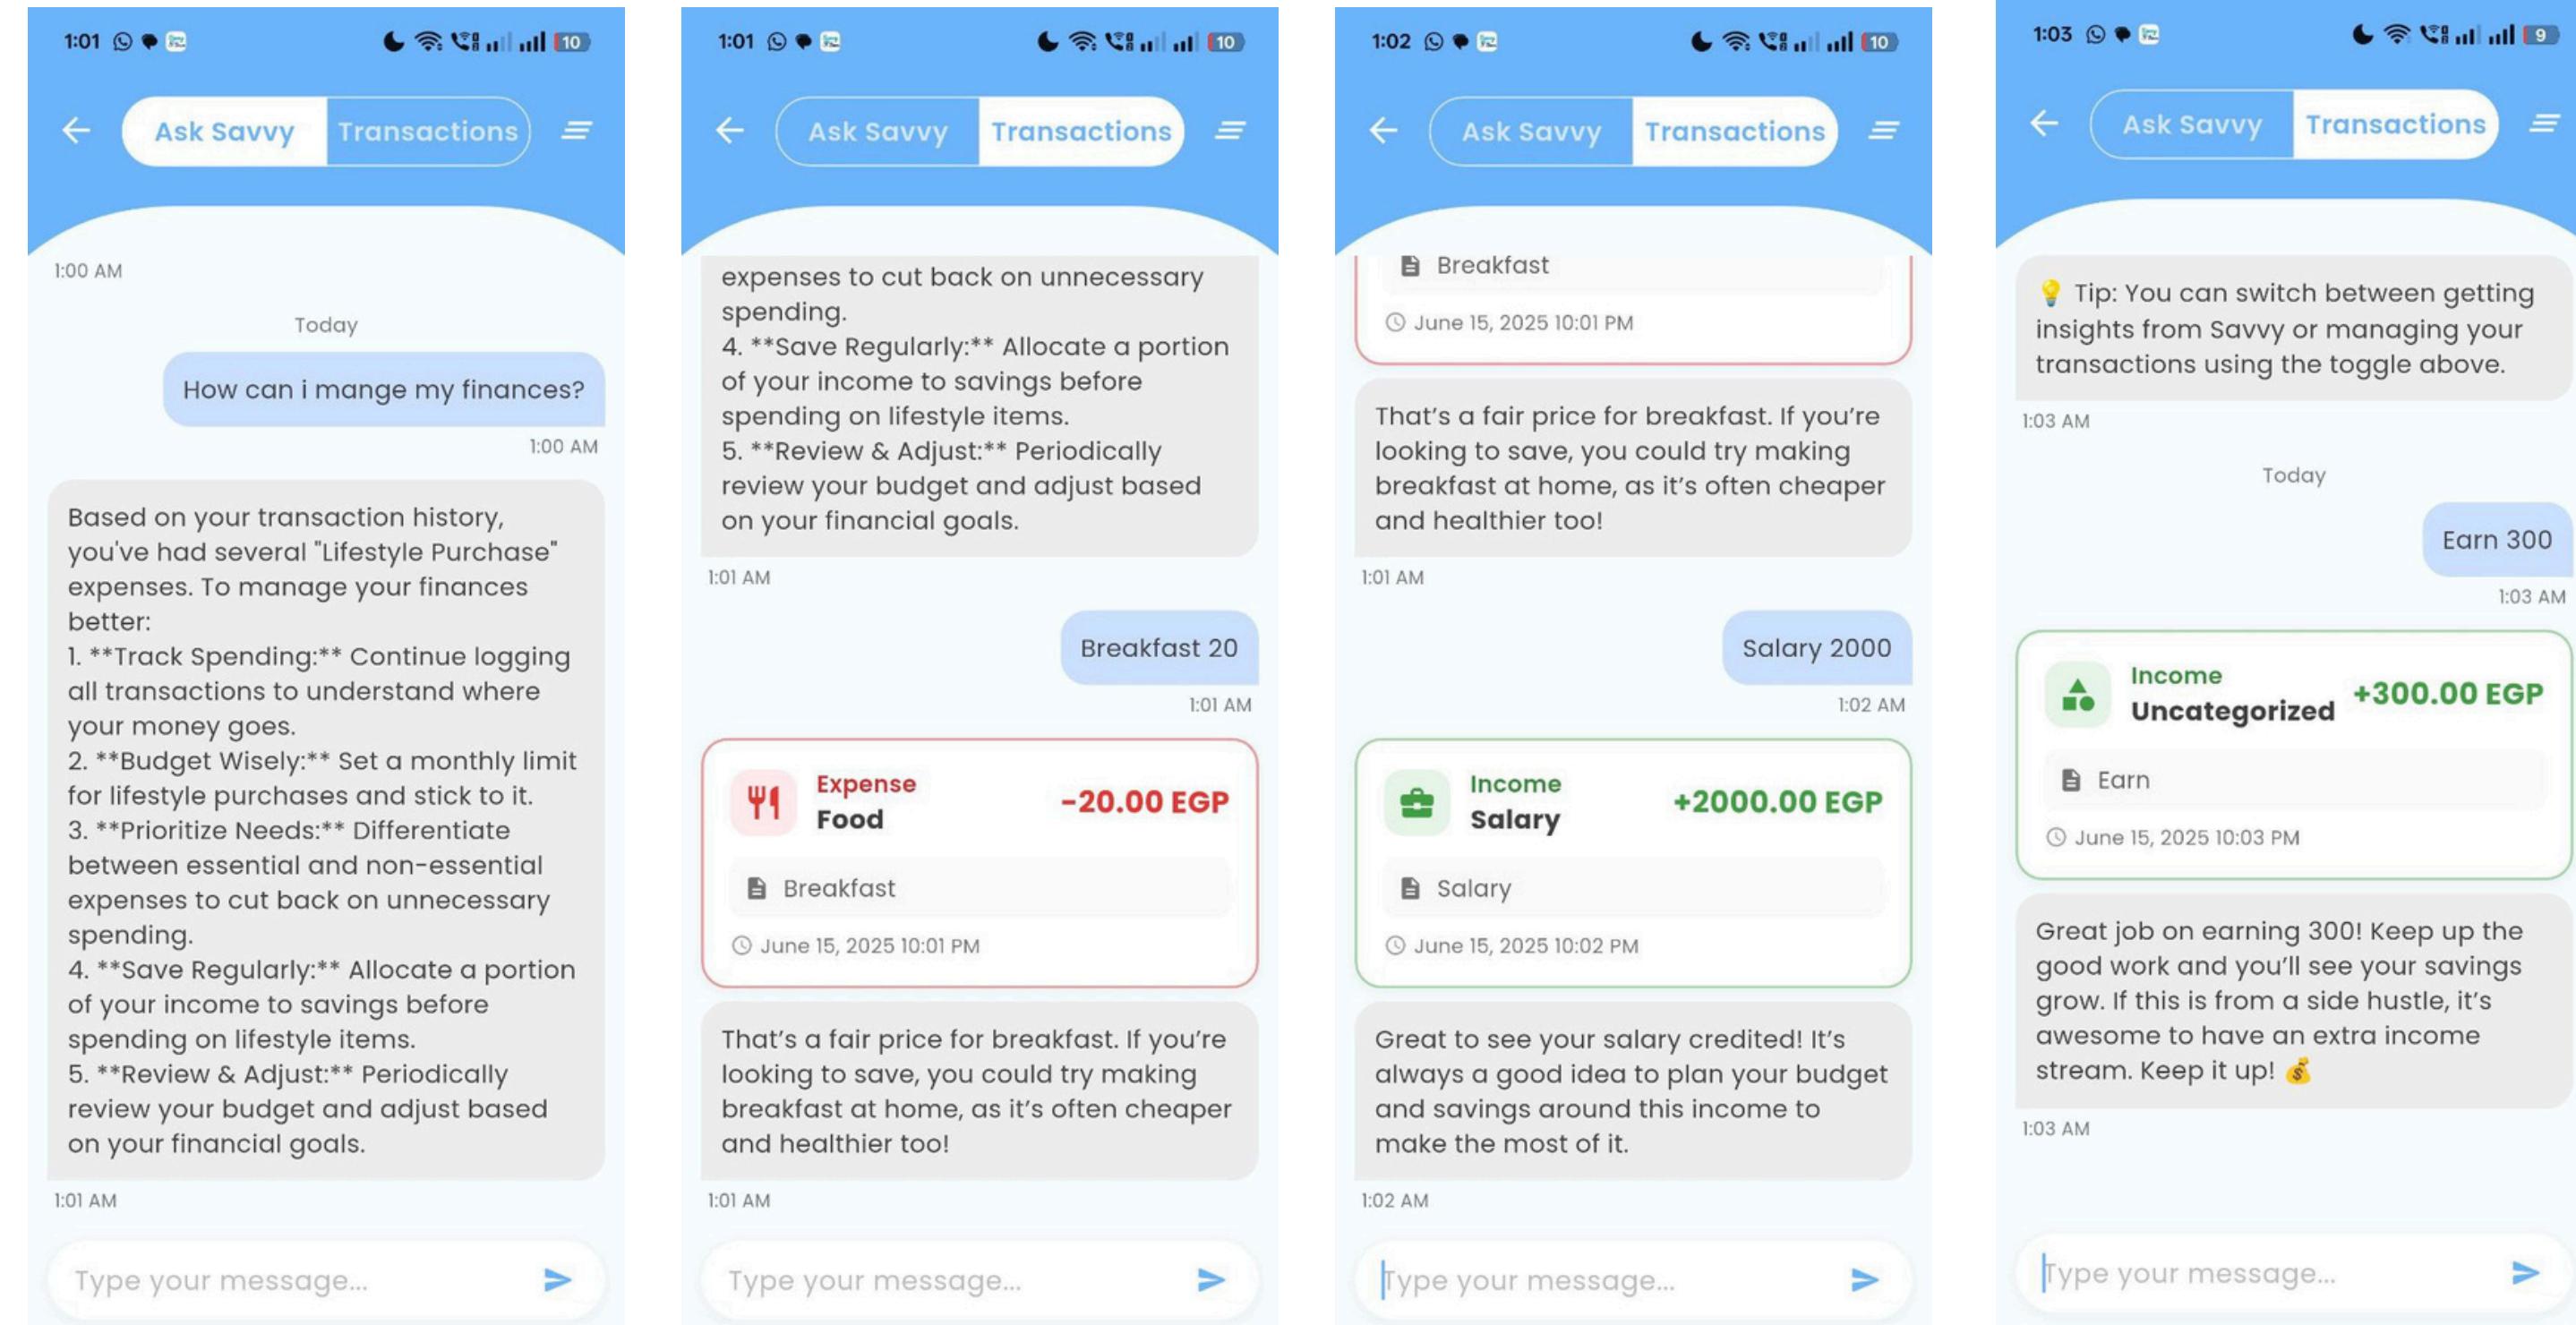
# Notifications



# Home screen

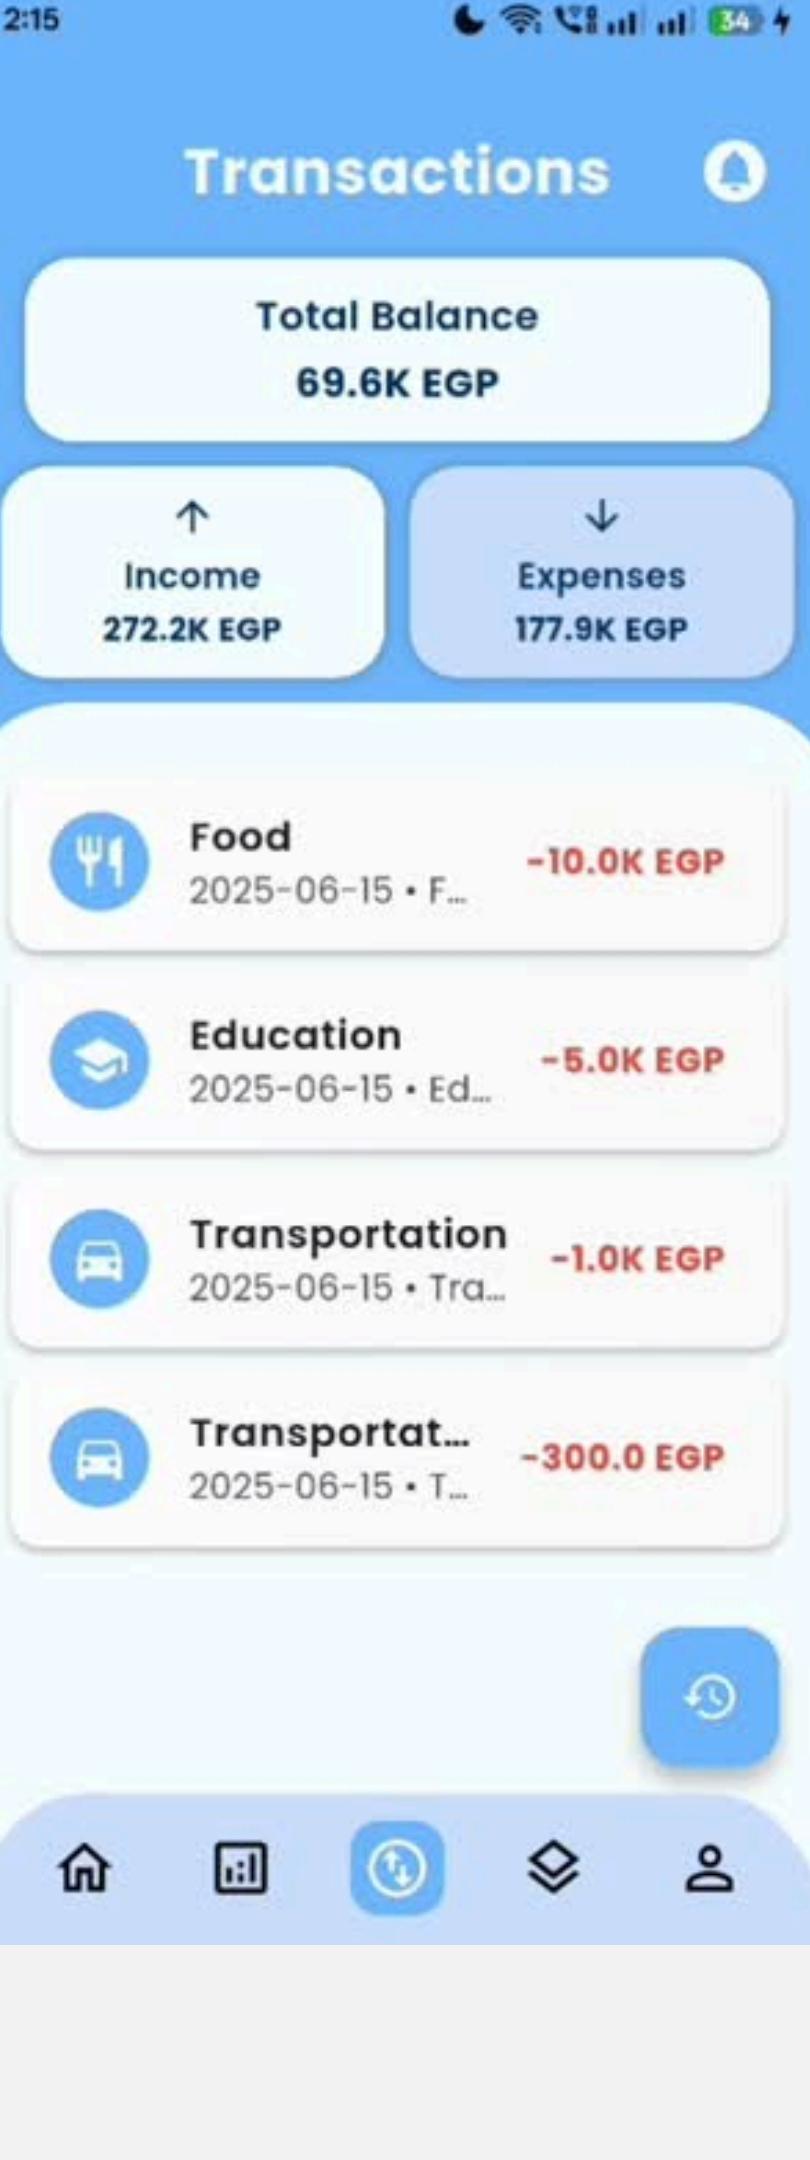


# chatbot screen



# Demo Time





# Time Series Analysis



# Data selection

**Goal:** Predict future monthly expenses to help users plan their budgets.

**Dataset Overview:** Contains:

1. **Date:** (from 9-2022 to 1-2025)
2. **Amount:** Expense amount.
3. **Category:** food, Transport,...etc.



**Benefits for Modeling:**

- Enables multivariate forecasting (use categories as features).
- Easy to preprocess: can group by month and category.
- Reflects personal budgeting behavior, aligned with app goals.

# Data preprocessing

- **Category Mapping Aligned with App Design.**
- **Outlier Detection & Removal (Ensemble Method)**
  - Combined KNN, Local Outlier Factor (LOF), and Isolation Forest to remove anomalous transactions.
  - Prevents skewing of predictions due to extremely large or rare transactions.
- **Monthly Aggregation**
  - Grouped daily transactions into monthly totals per app category.

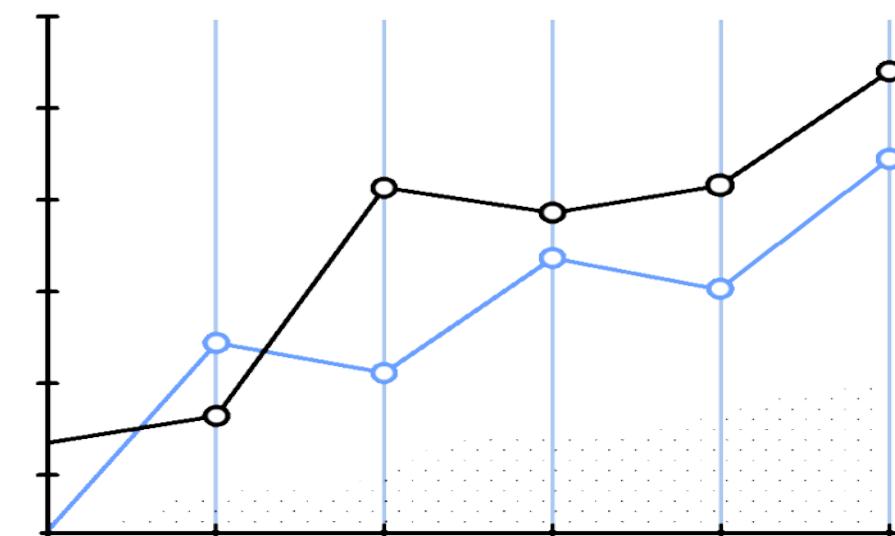


# Data Preparation

- **Features:**
  - "Education", "Entertainment", "Fashion", "Food", "Lifestyle", "Transportation", "Health"
- **Target:**
  - Expenses (total monthly spending).

## Preprocessing Steps:

- Standardization:
  - Applied Standard Scaler to all features.
  - Ensures consistent scale across different spending categories.
- Train/Test Split:
  - 85% of data for training.
  - The remaining 15% for testing.
  - No shuffling – preserves time order.



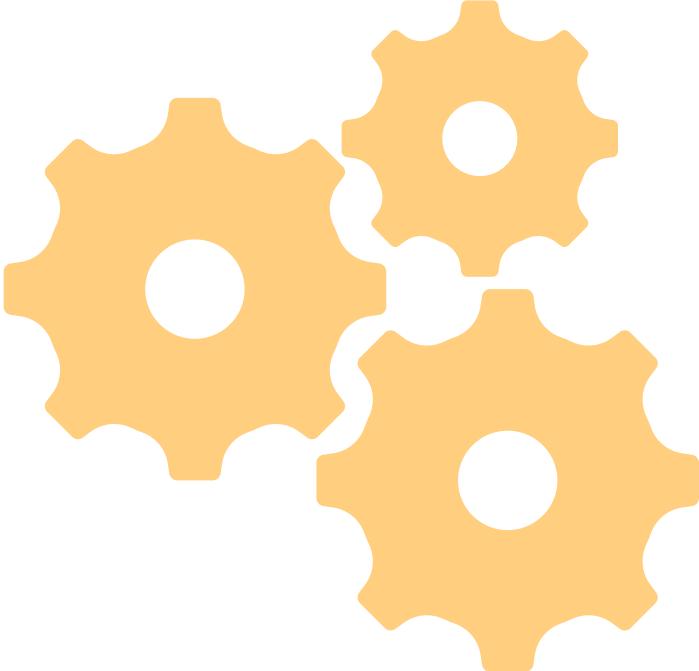
# Data Preparation

## **Sequence Design:**

- Sliding Window Size: 3 months
  - For each prediction, use the last 3 months of all 7 features.

Testing multiple window sizes (e.g., 1, 3, and 6) showed that 3 months consistently led to better predictive performance, as measured by lower error metrics (RMSE).

# Model Architecture

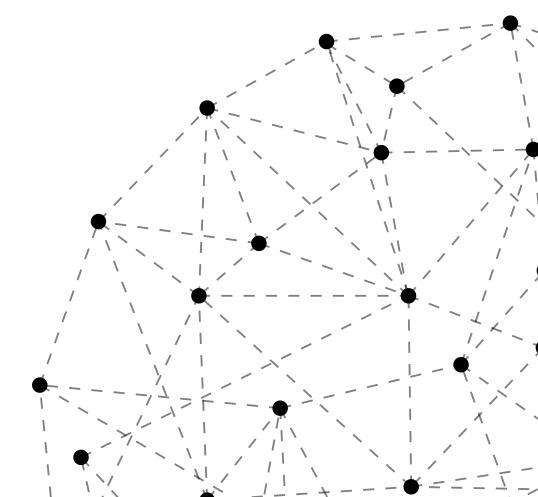


- **Input Layer**
  - Shape: 3 time steps × 7 features.
  - Each sample represents 3 months of spending across 7 categories.
  - Enables the model to learn temporal dependencies in multivariate data.
- **LSTM Layer 1 – 64 Units**
  - Processes the full 3-month sequence.
  - Learns short-term temporal patterns in each spending category.
  - Passes a sequence of hidden states to the next layer.
- **LSTM Layer 2 – 64 Units**
  - Takes output from LSTM 1 and condenses it into a single vector.
  - Captures higher-level sequential patterns across all categories.
  - Outputs a fixed-length representation of the input sequence.



# Model Architecture

- **Dense Layer – 64 Neurons (ReLU Activation)**
  - Transforms the learned temporal features into a richer nonlinear space
  - Allows the model to understand complex relationships between categories and expenses
- **Dropout Layer – 20%**
  - Randomly disables 20% of neurons during training
  - Helps prevent overfitting by making the model less reliant on specific paths
- **Output Layer – 1 Neuron**
  - Outputs a single continuous value
  - Represents the predicted total expenses for the next month.



# Model Evaluation – Results

- **Test Setup:**

- Input: Past 3 months of spending data across 7 categories.
- Output: Predicted total Expenses for the next month.
- Test Size: Final 15% of the dataset.

- **Evaluation Metric:**

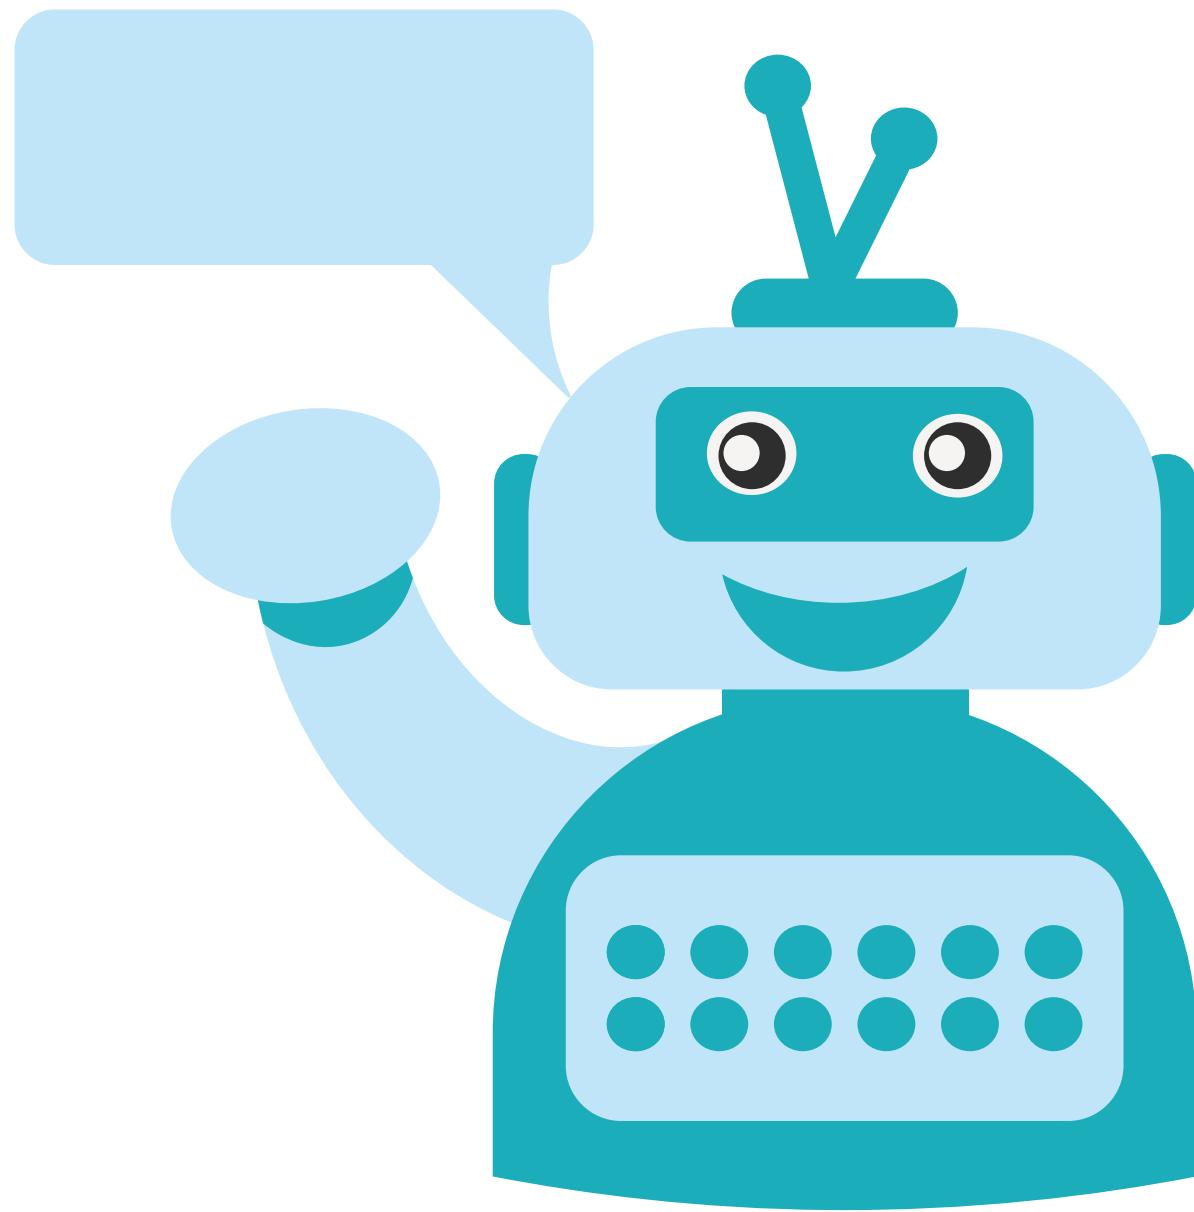
- Root Mean Squared Error (RMSE): 312.49



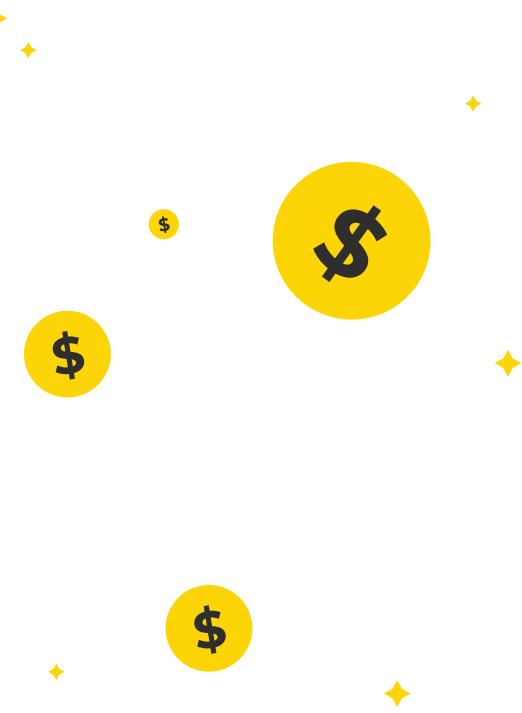
# Models Comparisons



Model	Input Type	RMSE
SARIMA	Univariate (past Expenses)	<b>840.15</b>
Univariate LSTM	Past 3 months of Expenses	<b>510.90</b>
Multivariate LSTM	Past 3 months Expenses of 7 Categories	<b>312.49</b>

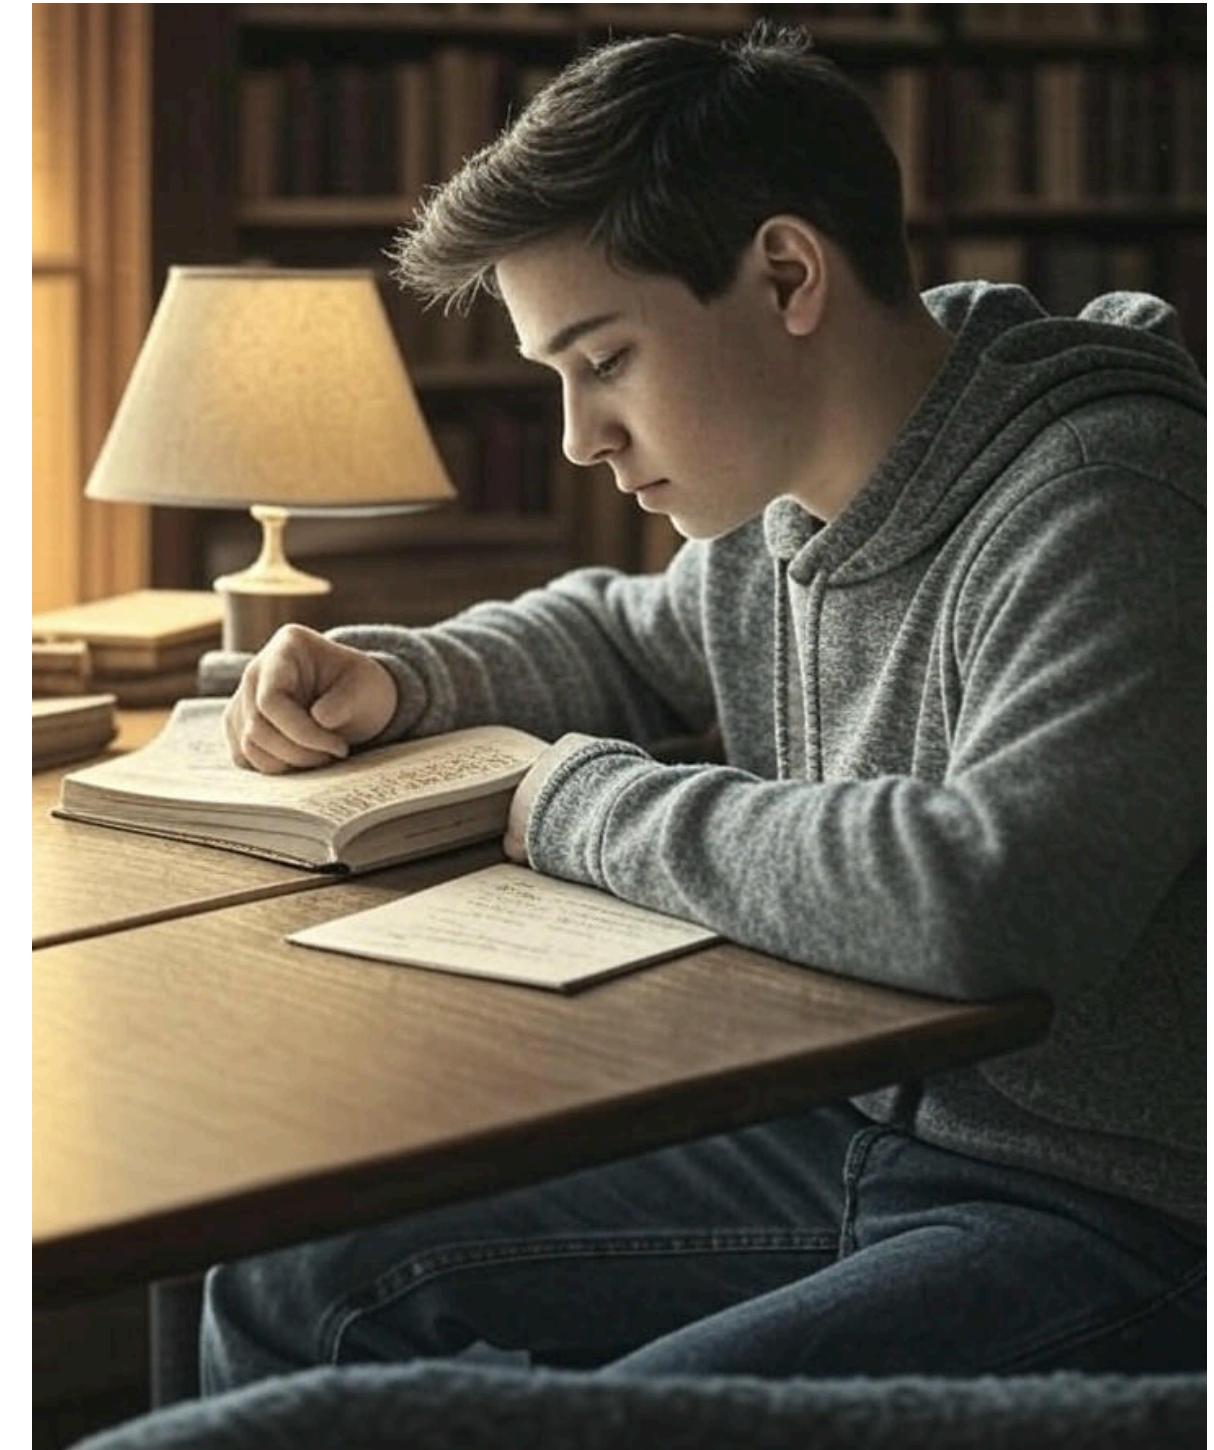


# AI-Chatbot



# What is RAG?

**Retrieval-Augmented Generation** is the process of optimizing the output of a large language model, so it references an authoritative knowledge base outside of its training data sources before generating a response.



# RAG Architecture Overview

## Step 1 : Documents Loading

Load data from various sources, including PDFs, text files, Web Pages, databases, CSV, JSON, Unstructured data, Research papers, and more.

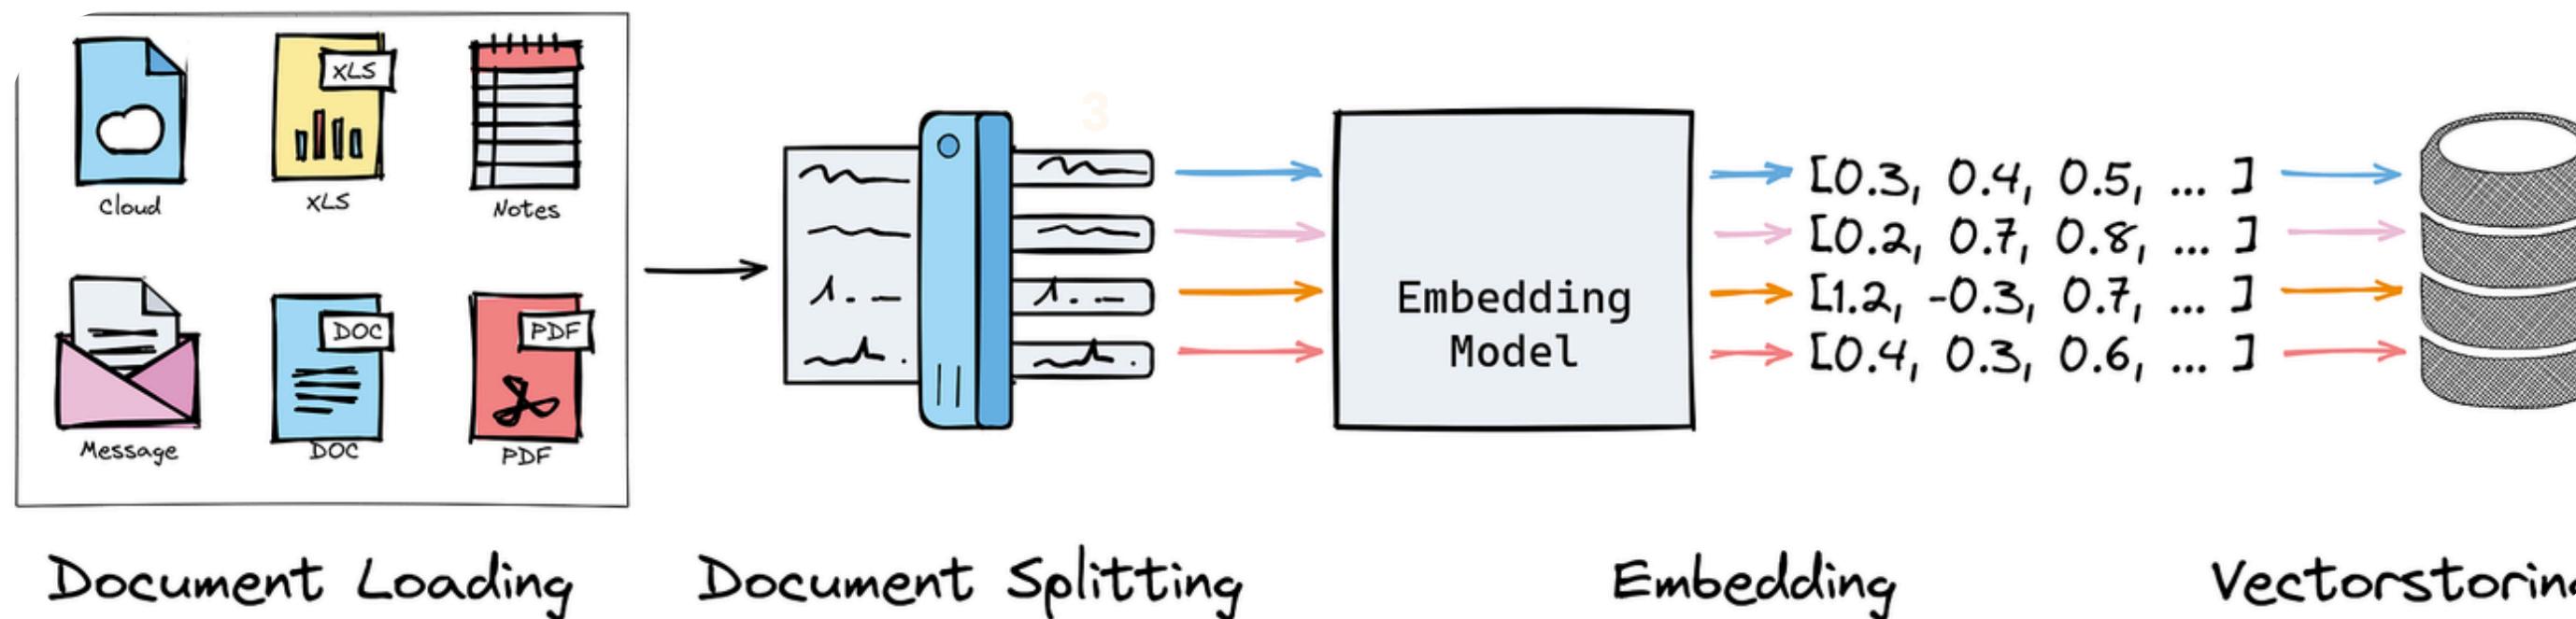
## Step 2 : Embedding generation

When the documents are loaded successfully, they are broken down into smaller parts or chunks and enable further processing by the language models.



## Step 3 : Vector storing

- Vector stores are specialized data stores that enable indexing and retrieving information based on vector representations.
- Vector stores are frequently used to search over unstructured data, such as text, images, and audio, to retrieve relevant information based on semantic similarity rather than exact keyword matches.

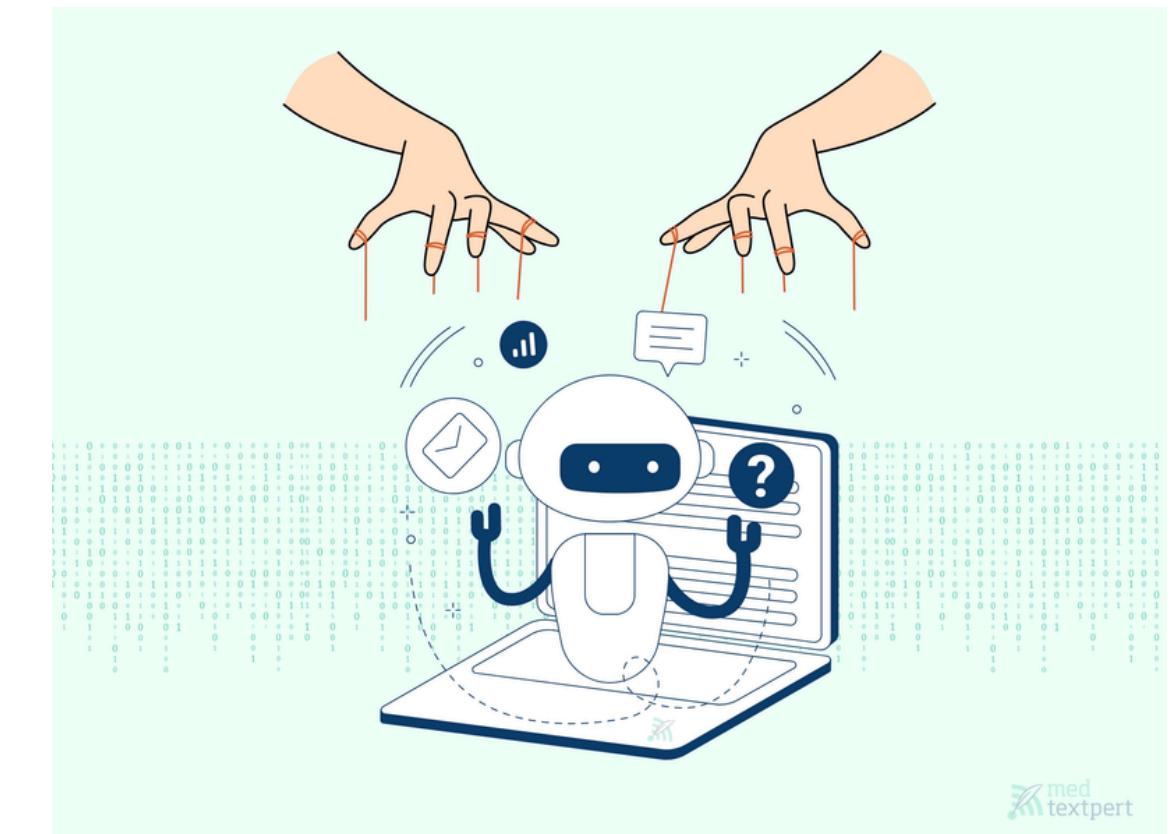


## Step 4 :Semantic Retrievers

- Semantic Retrievers focus on understanding the underlying context of a query and documents in order to retrieve the relevant information from the database.

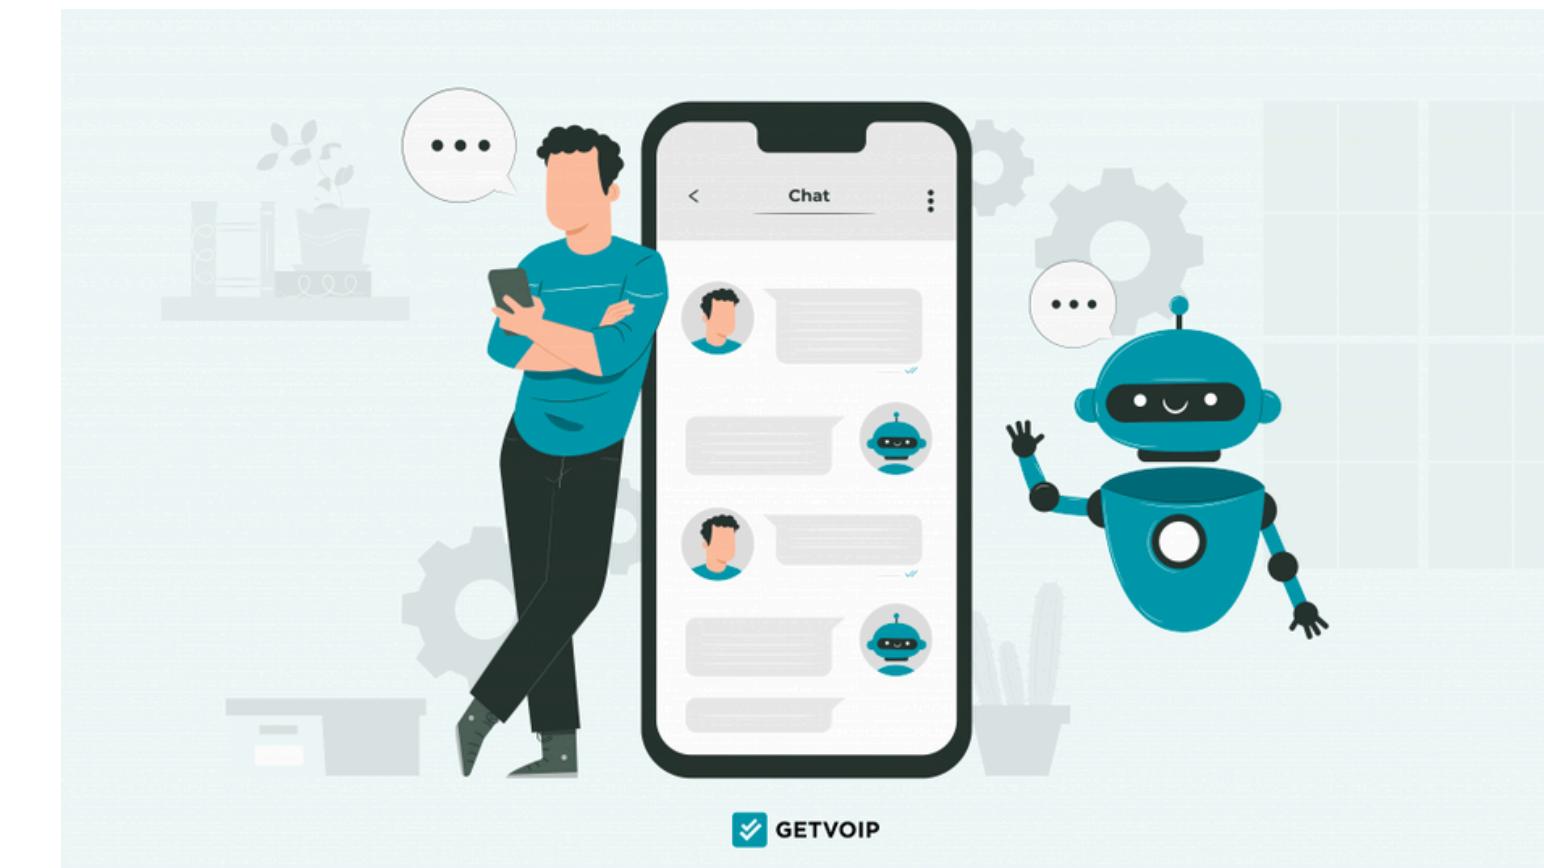
## Step 5 :Prompt Construction

- Providing a clear set of instructions and examples for LM's responses.



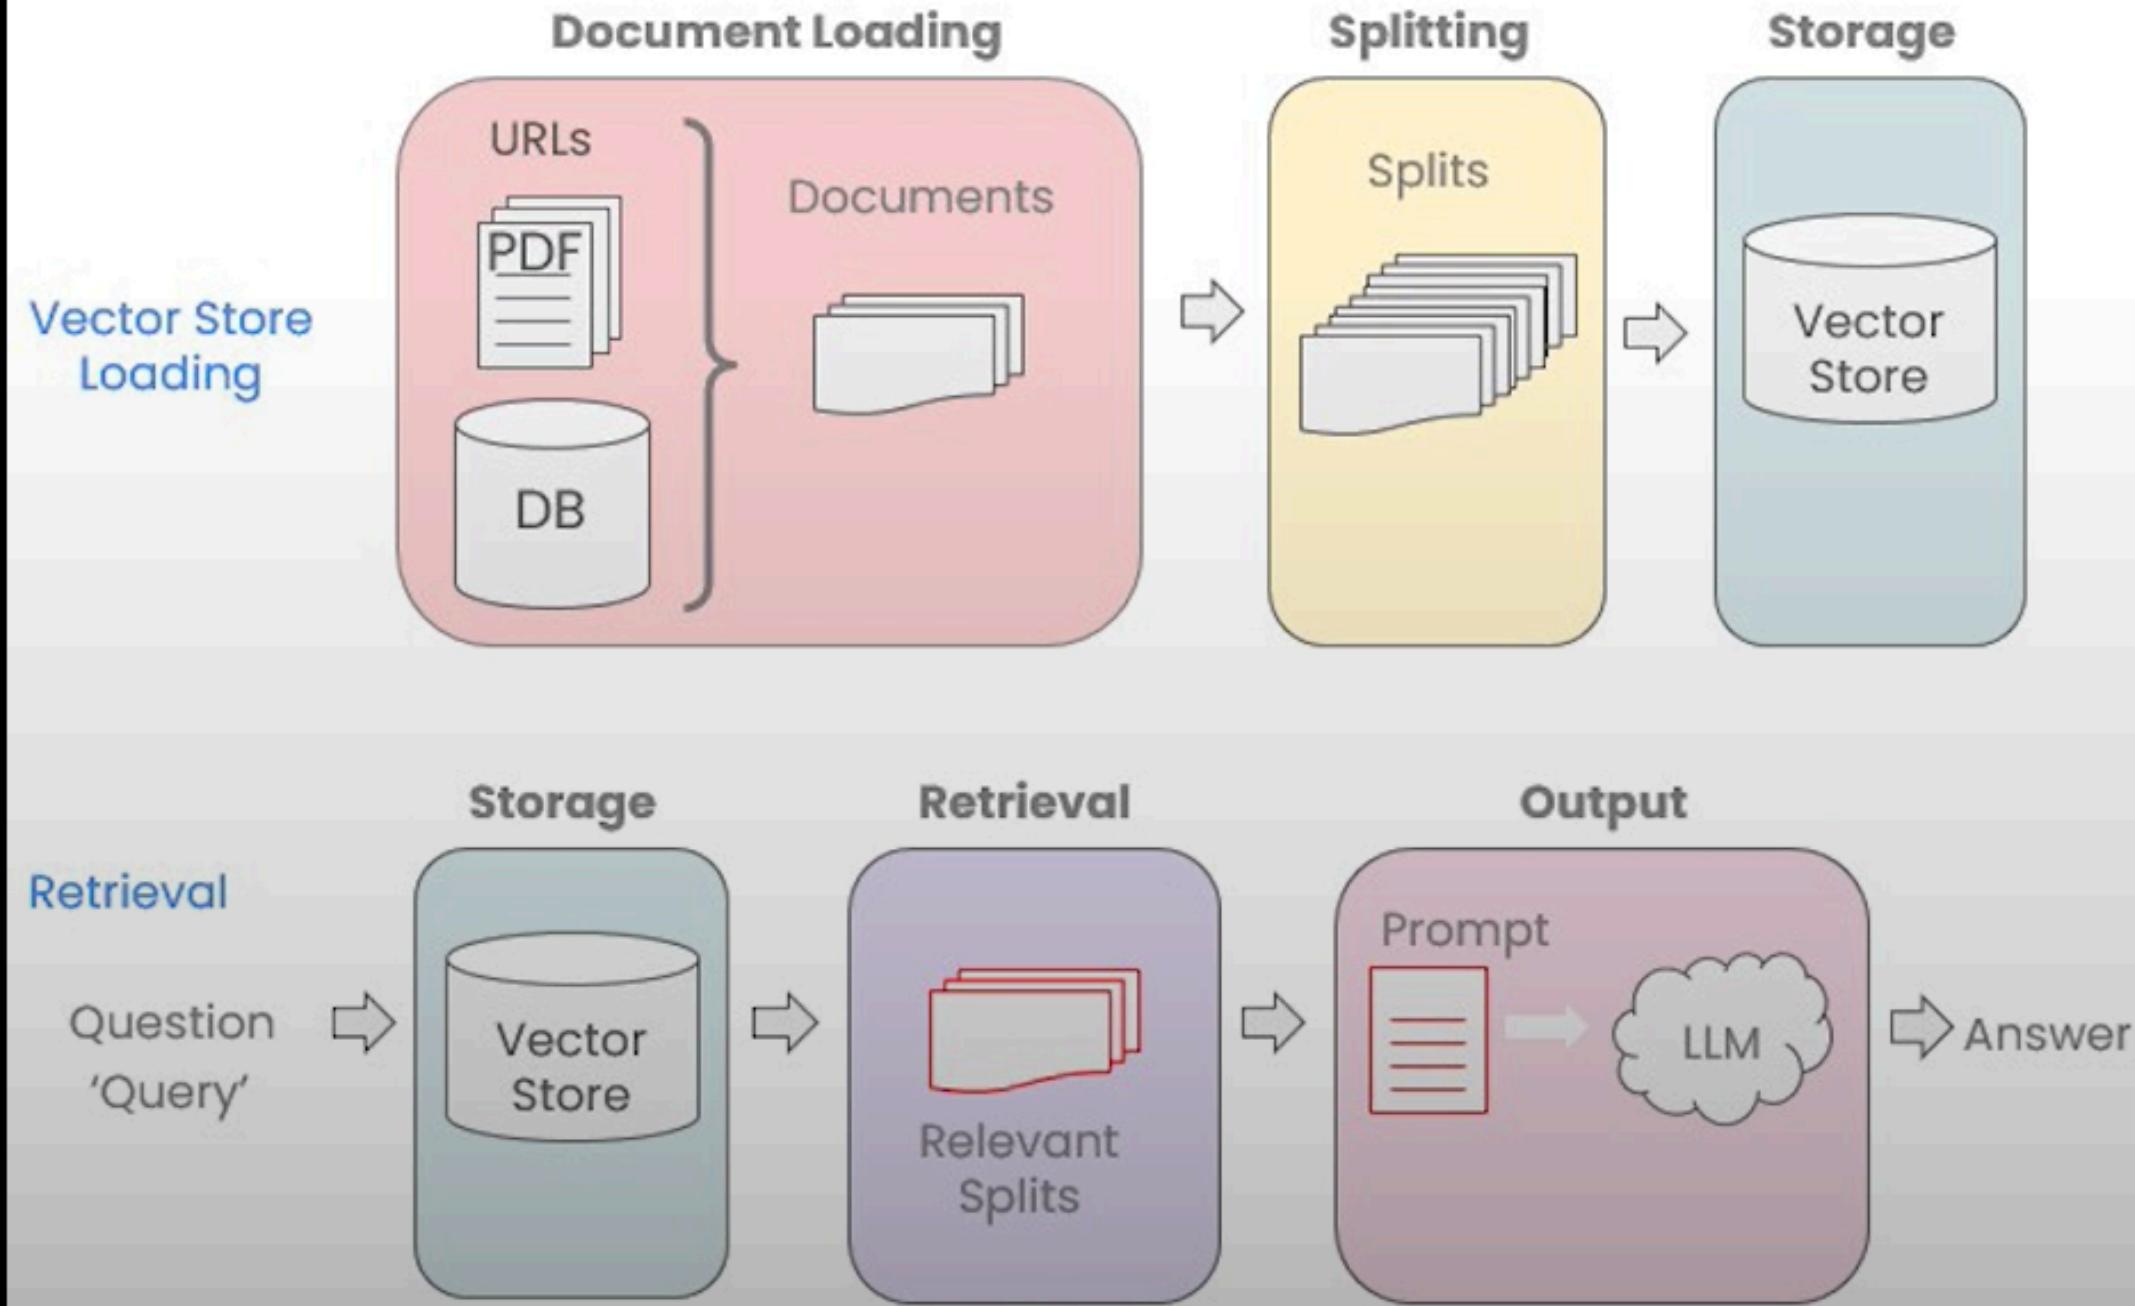
## Step 6 : Text generation

- Transformer model produces responses based on context.
- Deliver accurate, context-rich generated text.



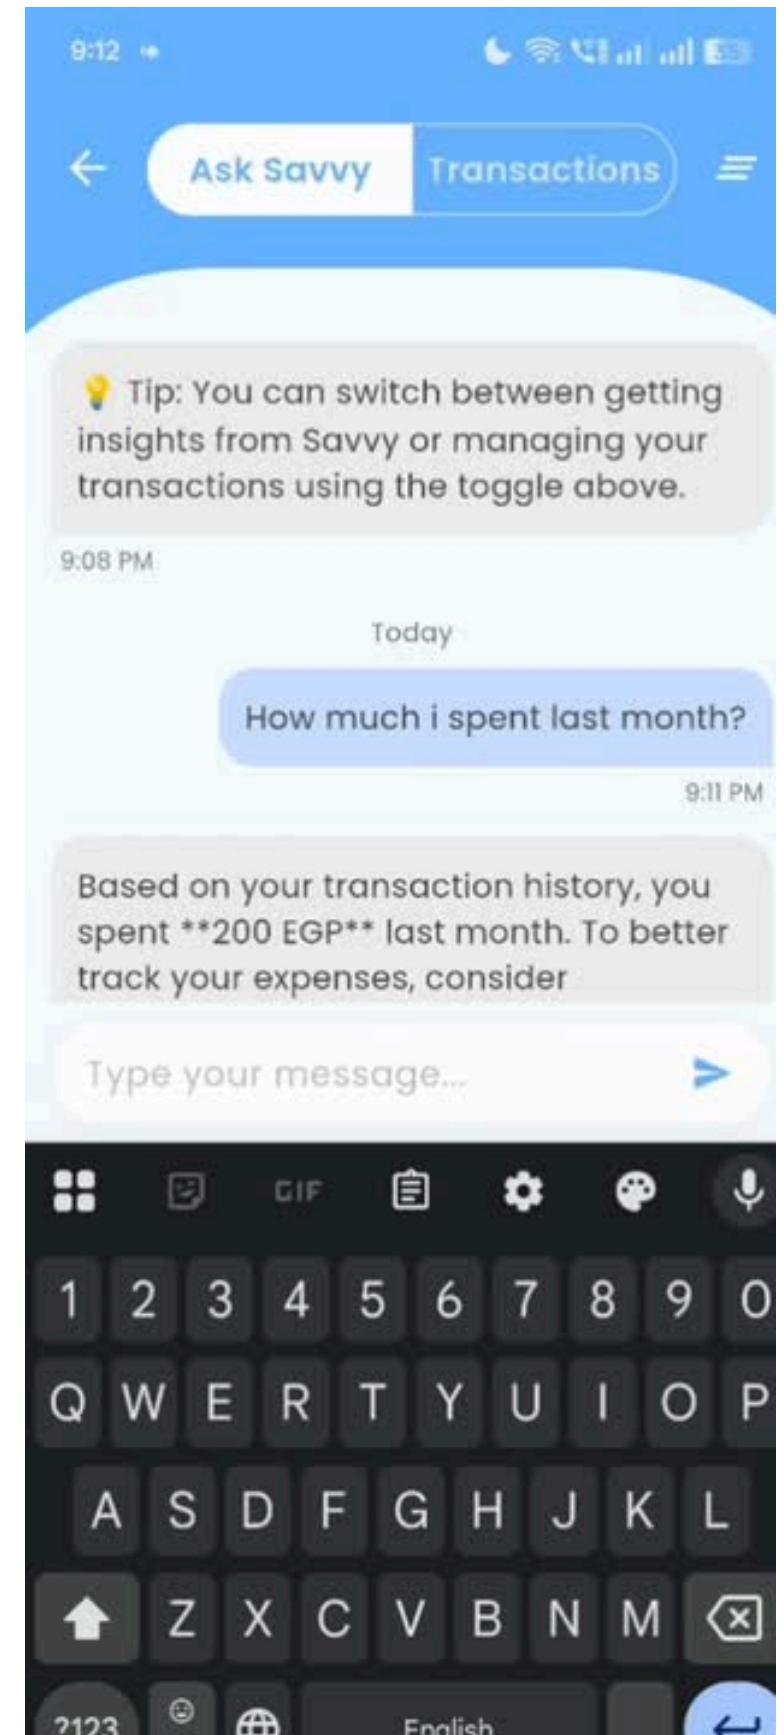
# LangChain

## Chat with Your Data



# Demo Time





# Why Transactions Chatbot?



**People forget to track daily expenses.**

**Amount**

\$ | 0.00

**Budgeting apps require manual input, boring and time-consuming.**

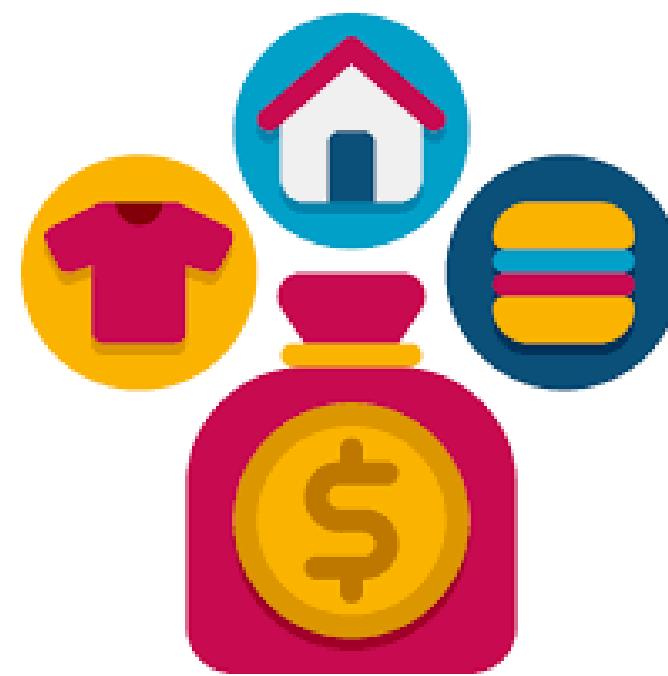


**Users crave real feedback, not just data storage.**

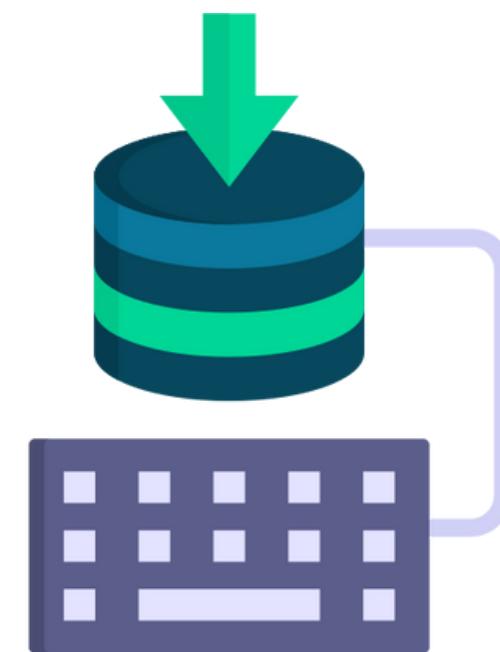
# Transactions Chatbot as a Solution



**Extract data  
from user input**



**Smart  
Categorization**



**Add transaction  
to Database**



**Provide  
Personalized  
feedback**

9:20

Add Expense

Date  
15/06/2025

Food

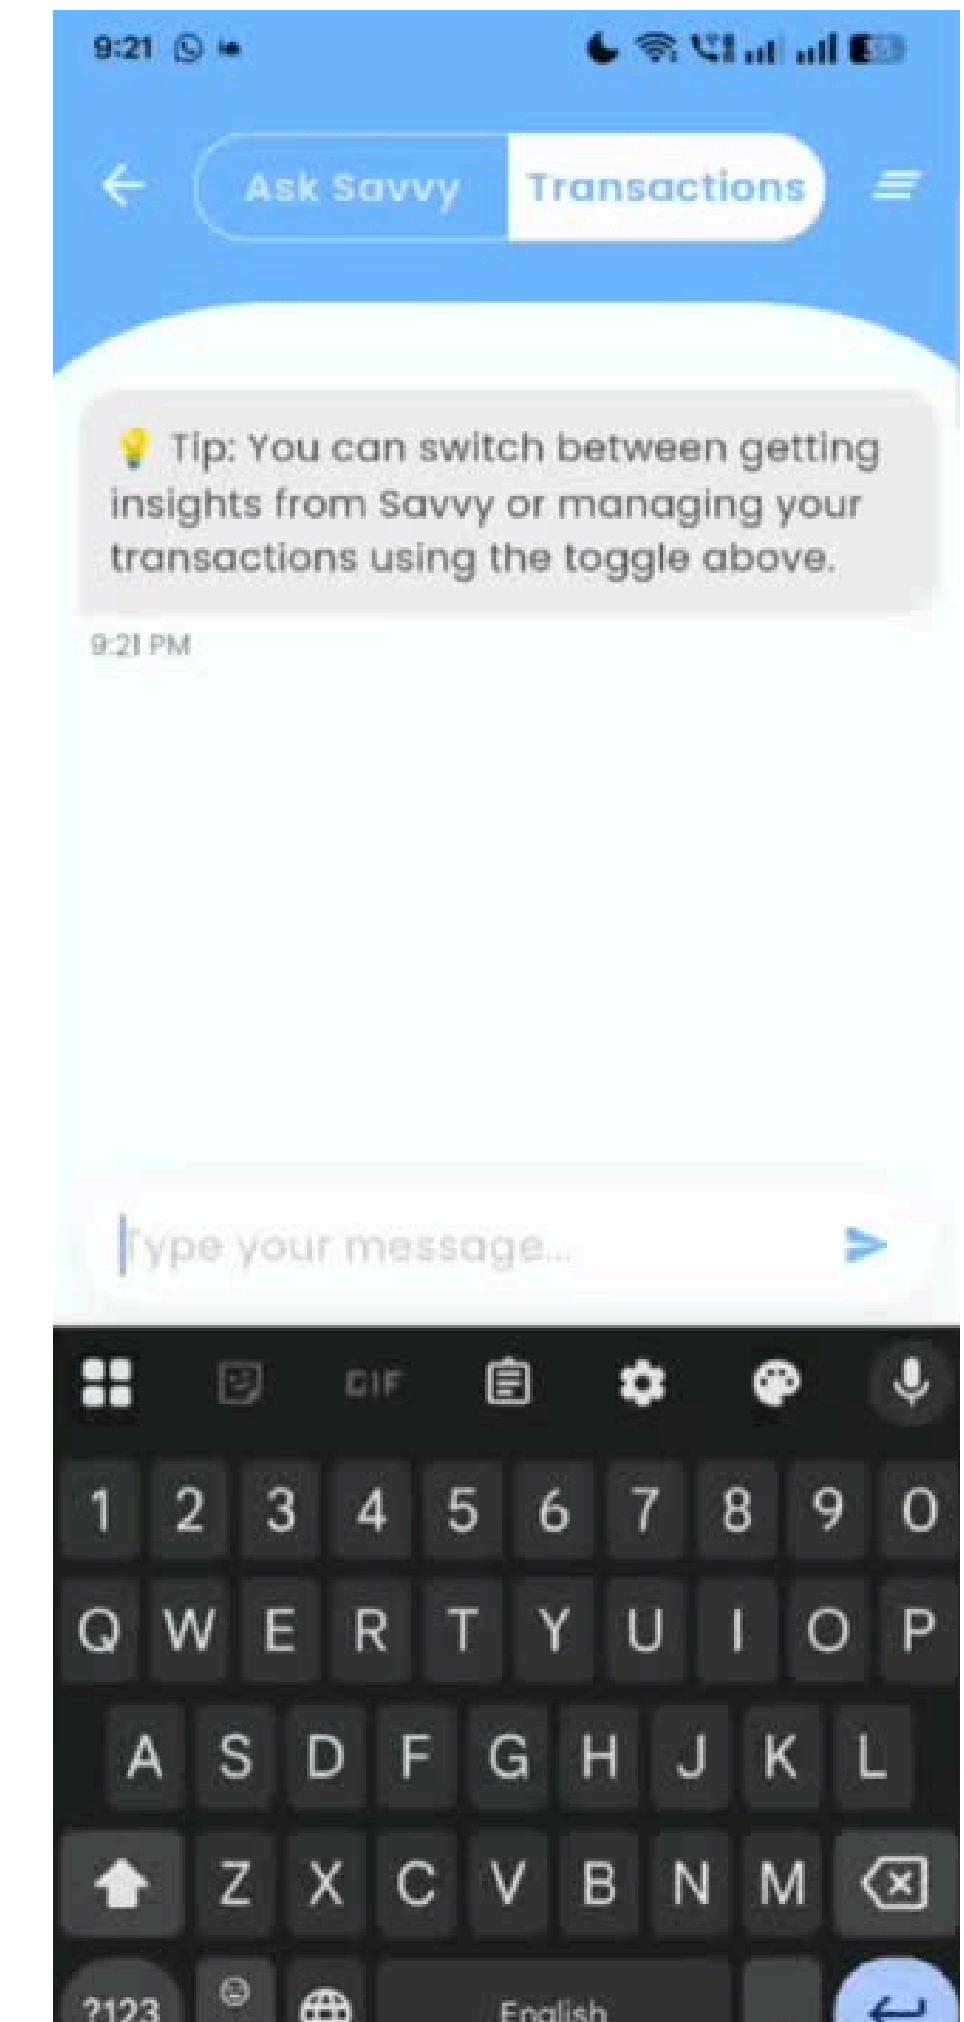
Amount  
EGP 00.00

Description (optional)

Save

Microphone, Camera, Settings, GIF, Image, Grid icons

This screenshot shows the 'Add Expense' screen of a mobile application. At the top, it displays the time (9:20) and battery level. Below that is the title 'Add Expense'. The form consists of several input fields: 'Date' set to '15/06/2025', a dropdown menu currently showing 'Food', 'Amount' set to 'EGP 00.00', and an optional 'Description' field. A large blue 'Save' button is at the bottom. A black navigation bar at the very bottom contains icons for microphone, camera, settings, GIF, image, and a grid.





# The Journey

## BERT-based System (Phase 1)

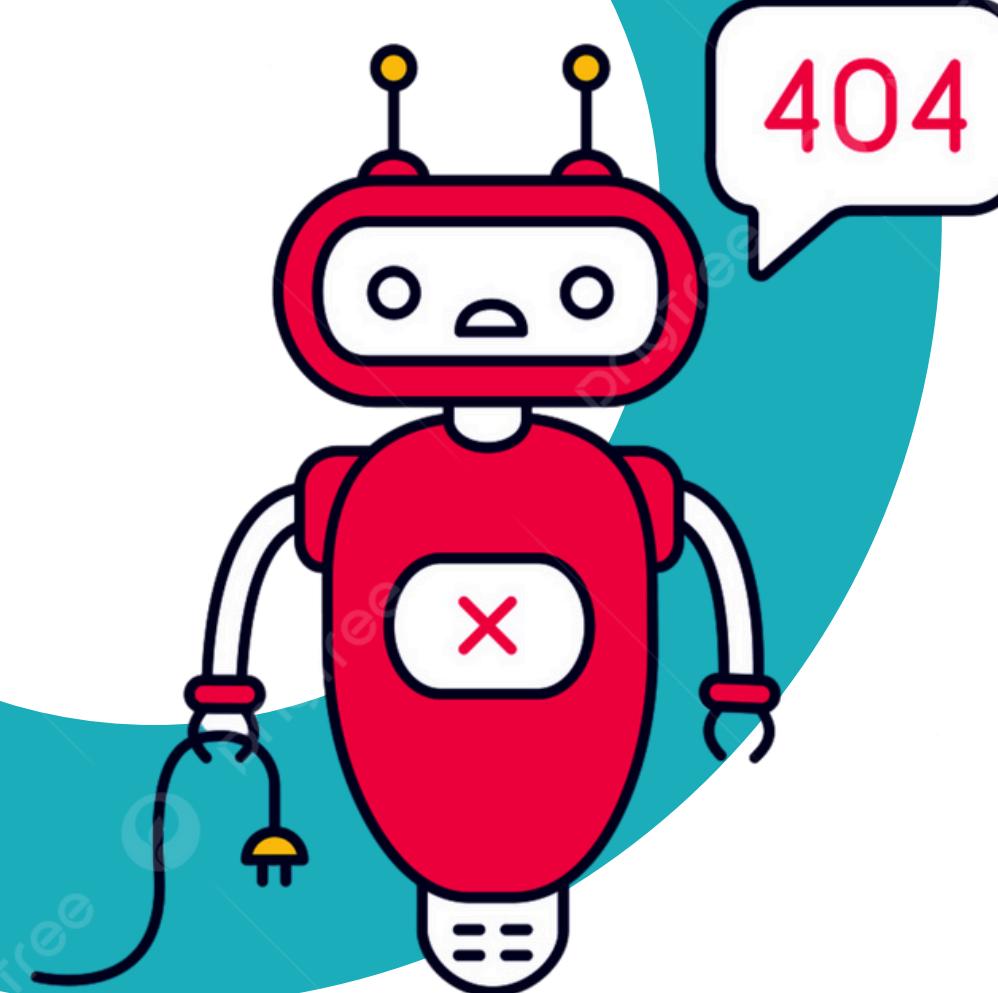
- Limited NLU
- Decent Classification
- Robotic & bland feedback

## Hybrid Prompting (Phase 2)

- Intent collision between flows
- Fuzzy boundaries, unclear outputs

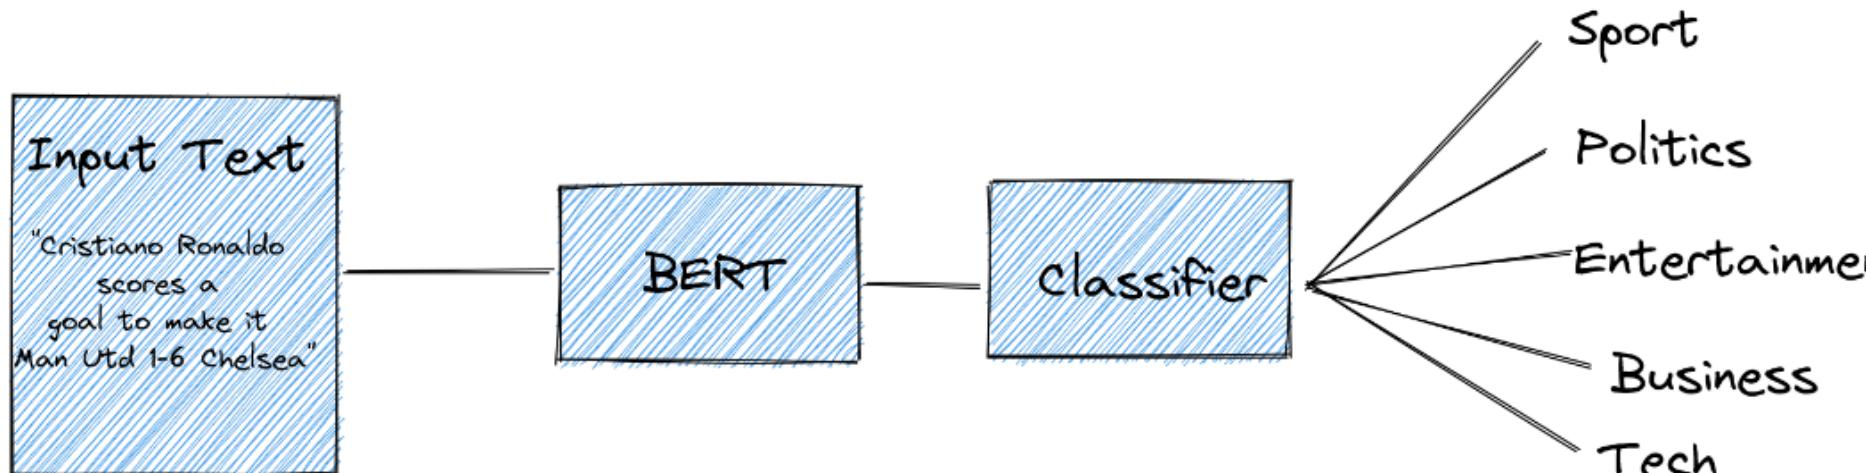
## Custom LLM Prompting (Phase 3)

- Conversational, context-aware, robust
- Seamless entity extraction, smart categorization
- Feedback that feels human

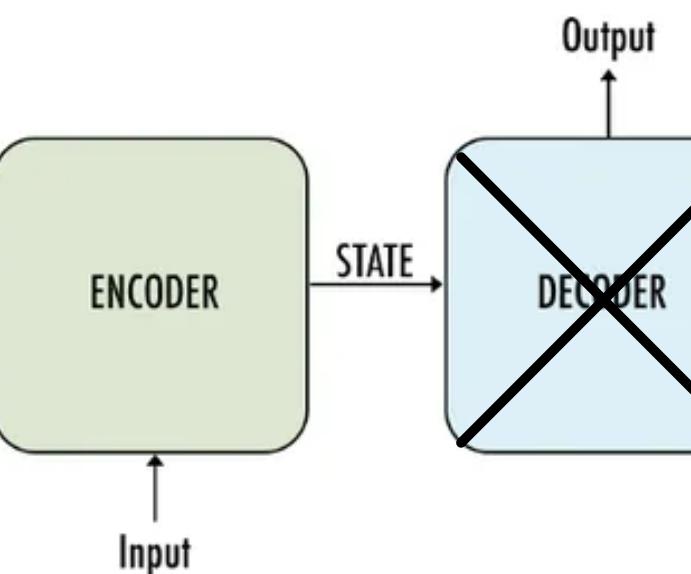


# From BERT to Breakthrough

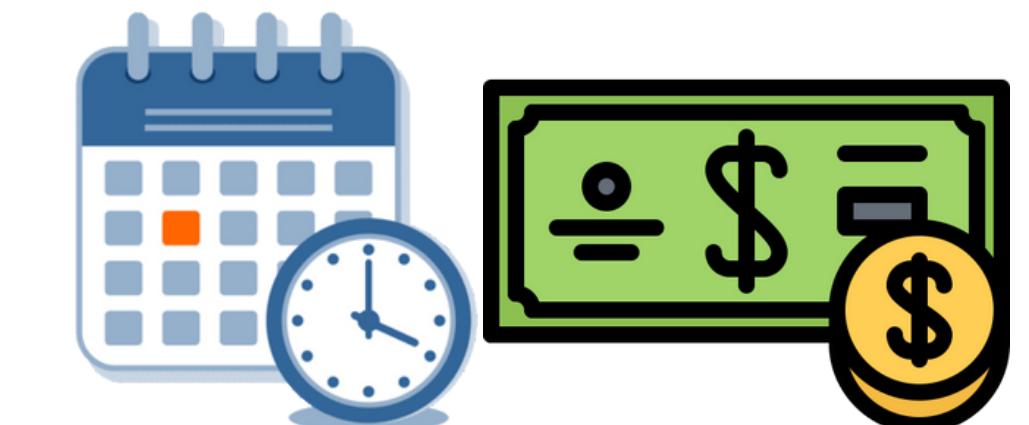
**Decent Classification**



**Encoder Model**



**Limited Natural language understanding**



Date and time

Amount



# Combined Prompt

## Intent Recognition

User Input — Process — Classify

I am looking for  
Mexican restaurant

1

Input Data

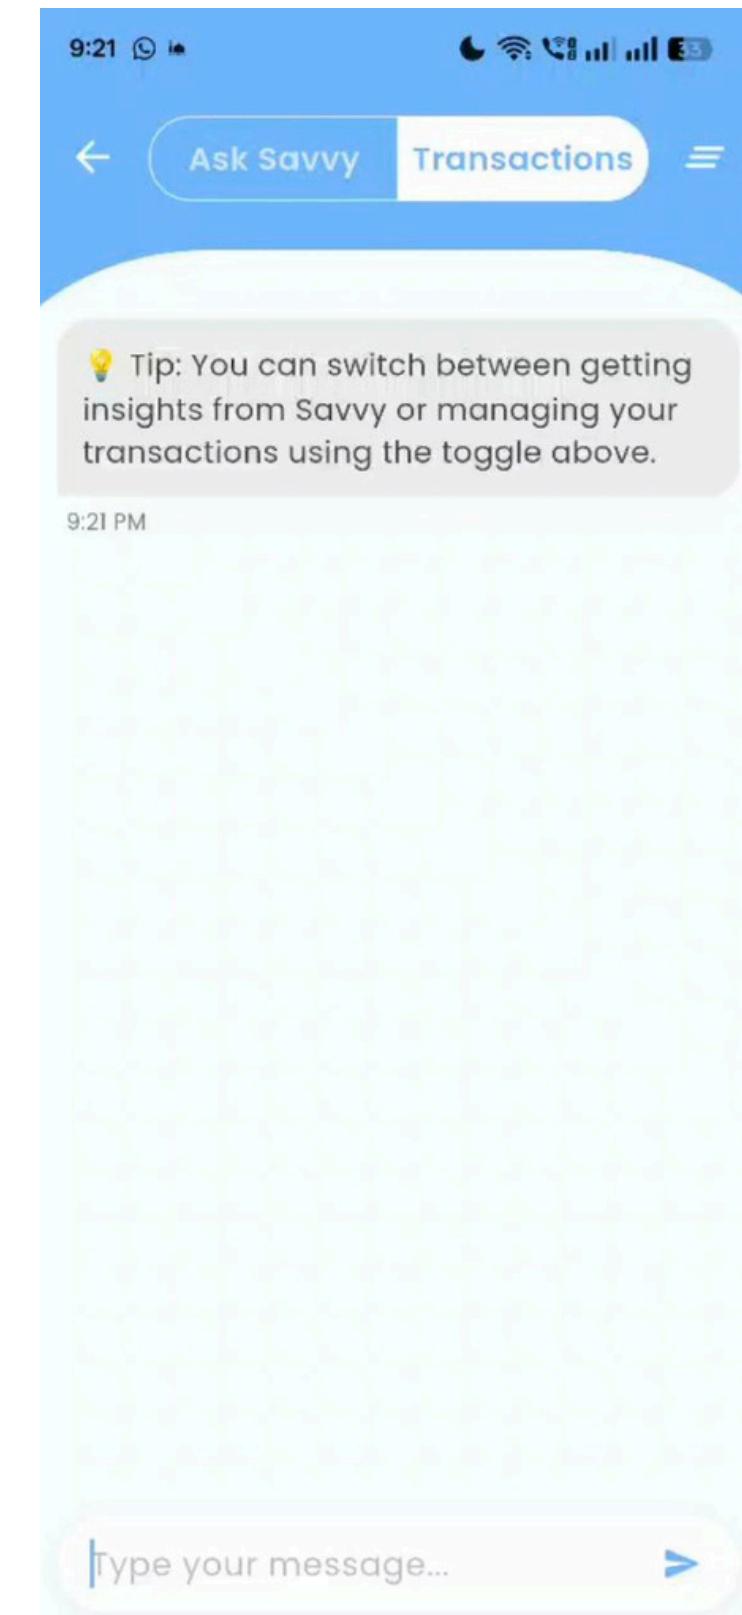
2

Intent  
Recognition

3

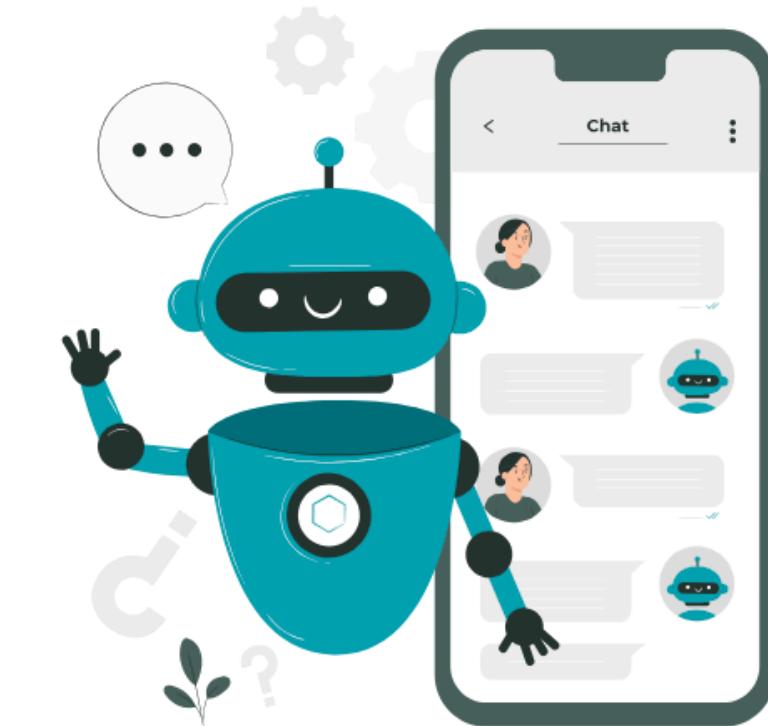
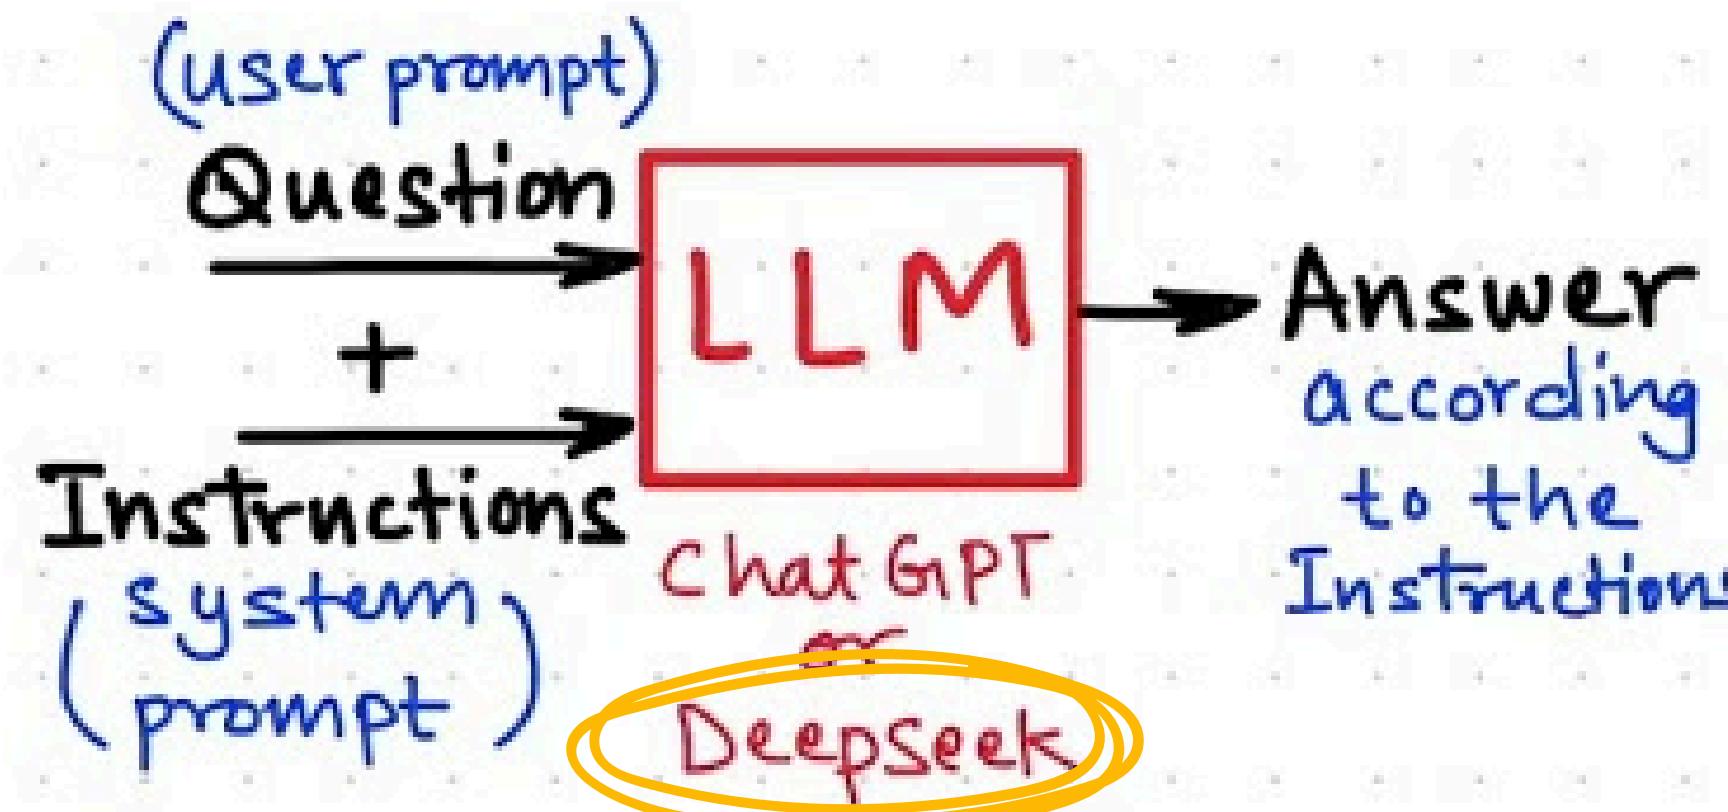
Response  
Output

# Separated Tasks



# The Breakthrough: Custom LLM Prompting

## [Prompt Engineering]





# Prompt Magic: Design Principles

1

## Structured Output

Enforced template (CREATED\_AT, AMOUNT, etc.)

2

## Natural Date Parsing

Understands "yesterday," "this morning," etc.

3

## Controlled Categories

Only allowed options, always consistent

4

## Conversational Feedback

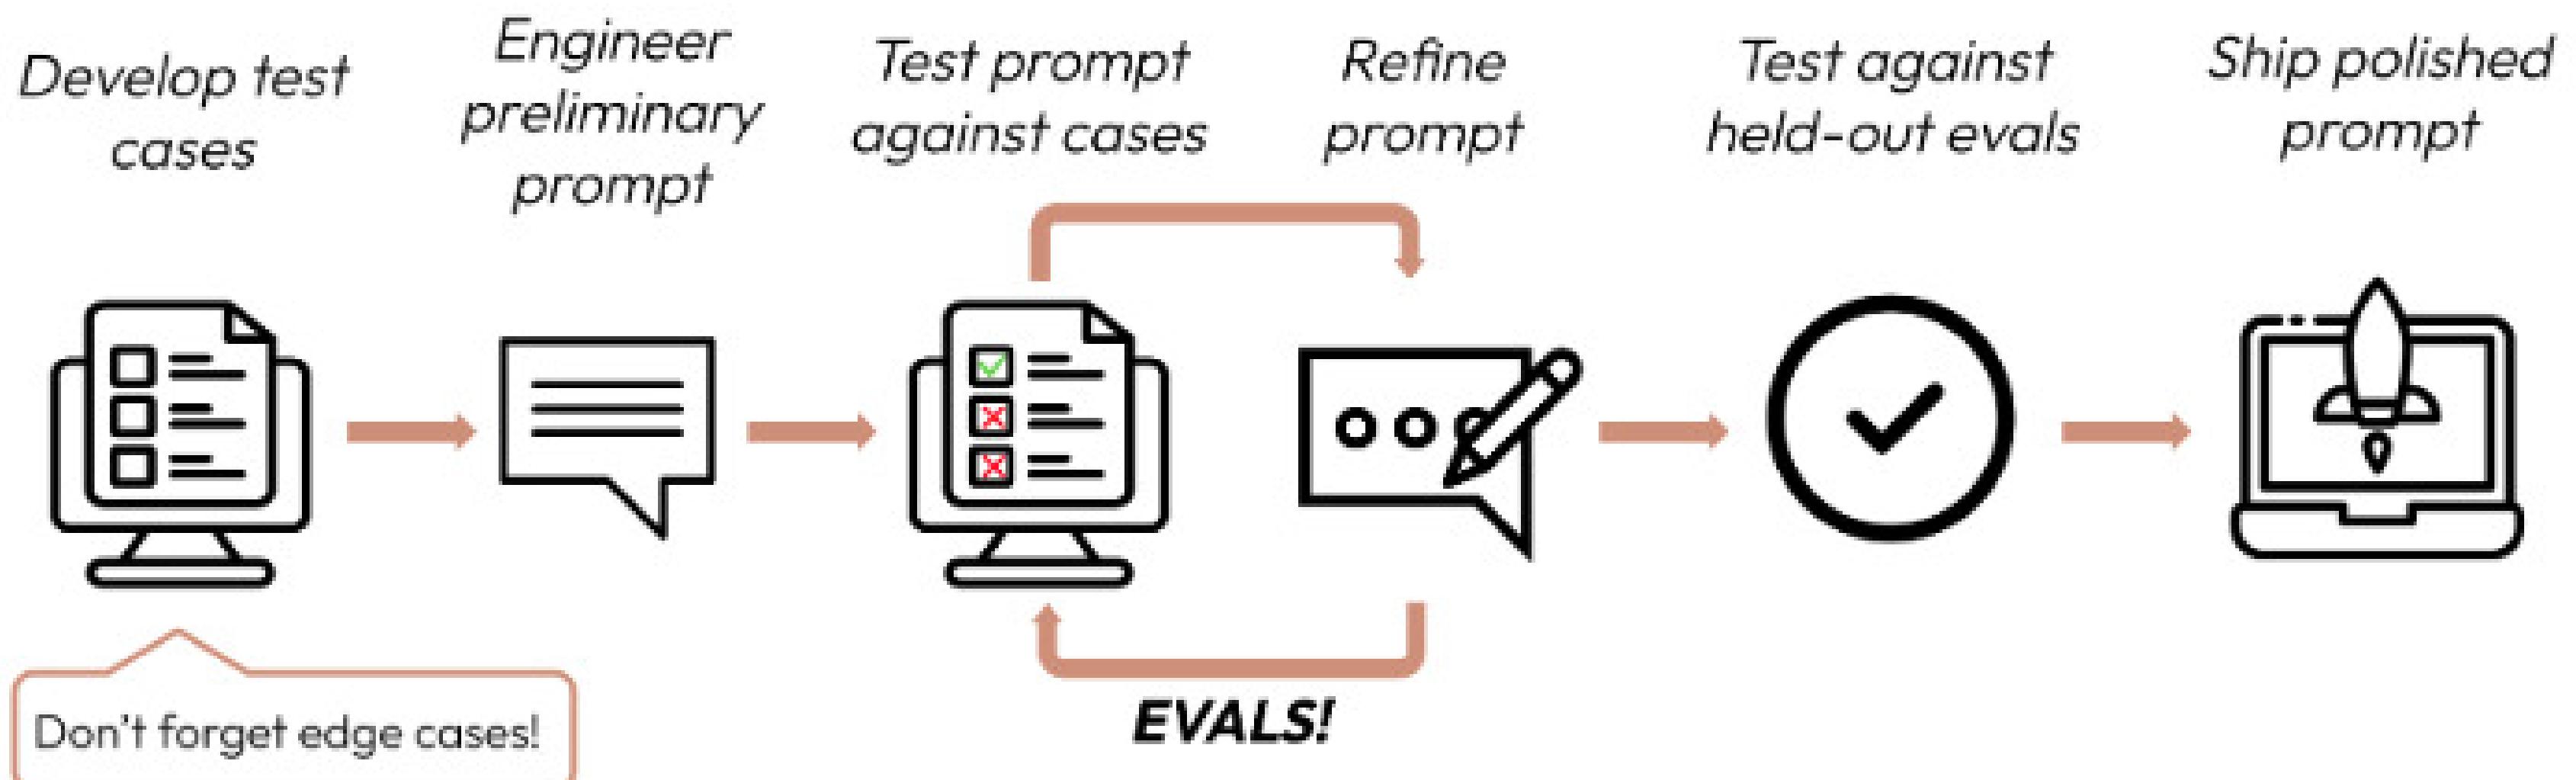
Empathetic, actionable advice (not just "OK!")

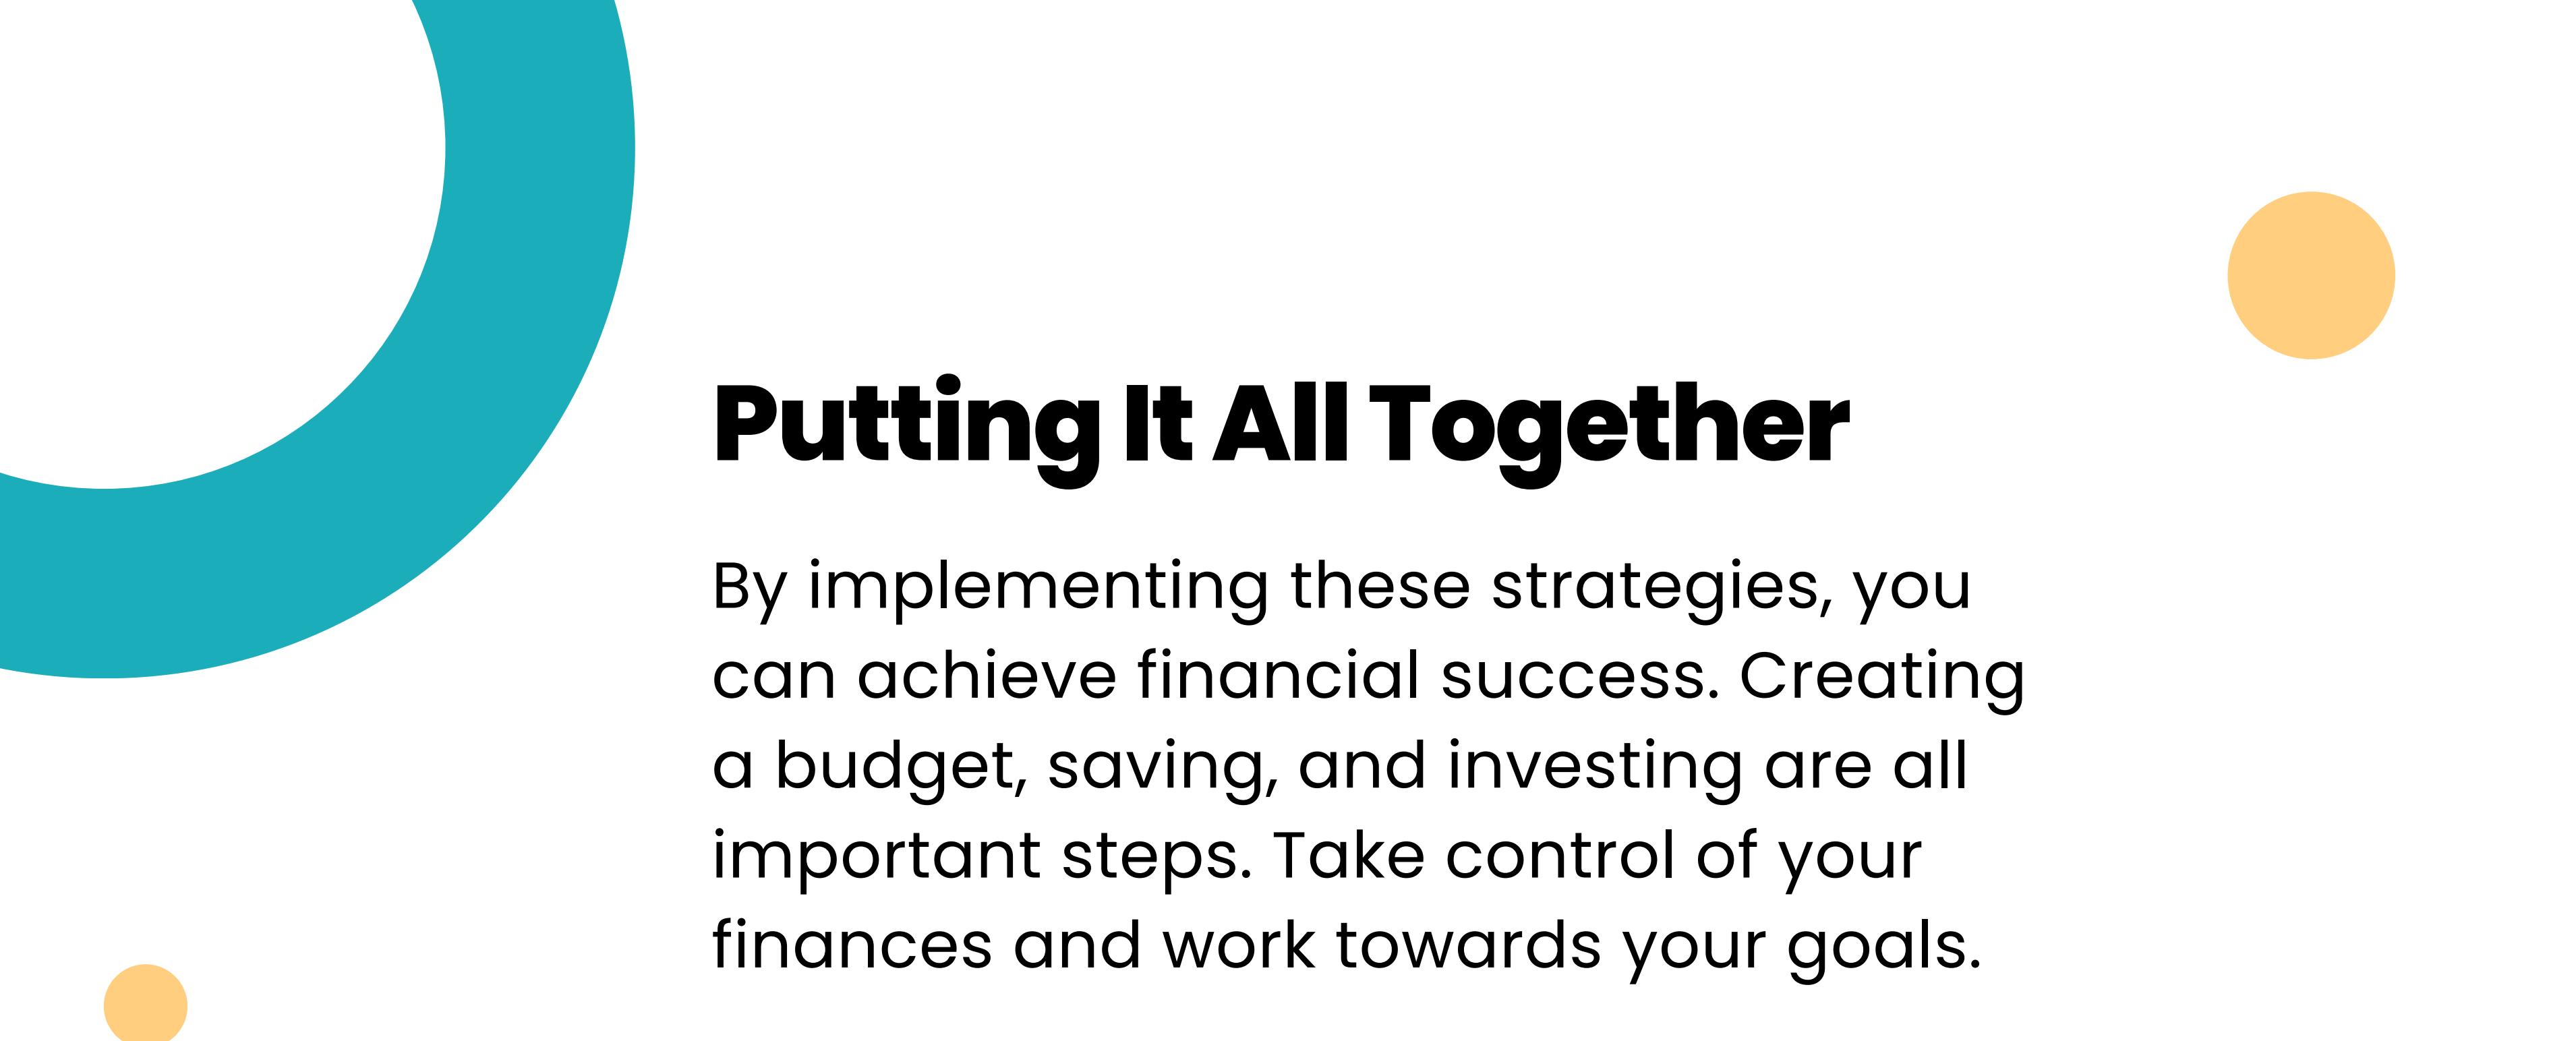
5

## Few-Shot Examples

Guides LLM for clarity & reliability

# Final Step: Test and Evaluation





# **Putting It All Together**

By implementing these strategies, you can achieve financial success. Creating a budget, saving, and investing are all important steps. Take control of your finances and work towards your goals.

# **Future Work**

**User can custom their own categories**

**Goals Improvements**

**Custom Date Range Filtering**

# Future Work

## Geofenced Notifications

app can notify users when they enter high-spending zones

(like shopping malls) with motivational or humorous alerts.



*Thank You*